

# KUBERNETES



Wojciech Barczyński - [Codility\\_](#)

# STORY

Kubernetes + Go + ...

- **Codility\_** - remote recruitment platform
- **SMACC/Hypatos** - Fintech / ML
- **Lyke** - Mobile Fashion app / ML

# Goal: continuous deployment

- Development: max 10 min  
Staging: max 15 min  
Production: max 20 min
- X to XX changes to production
- Infrastructure should be invisible

# Goal: continuous deployment

- Frequent releases
- Higher dev satisfaction
- Less (big) bugs
- Faster feature delivery

see: [Accelerate: The Science of Lean Software and DevOps](#)

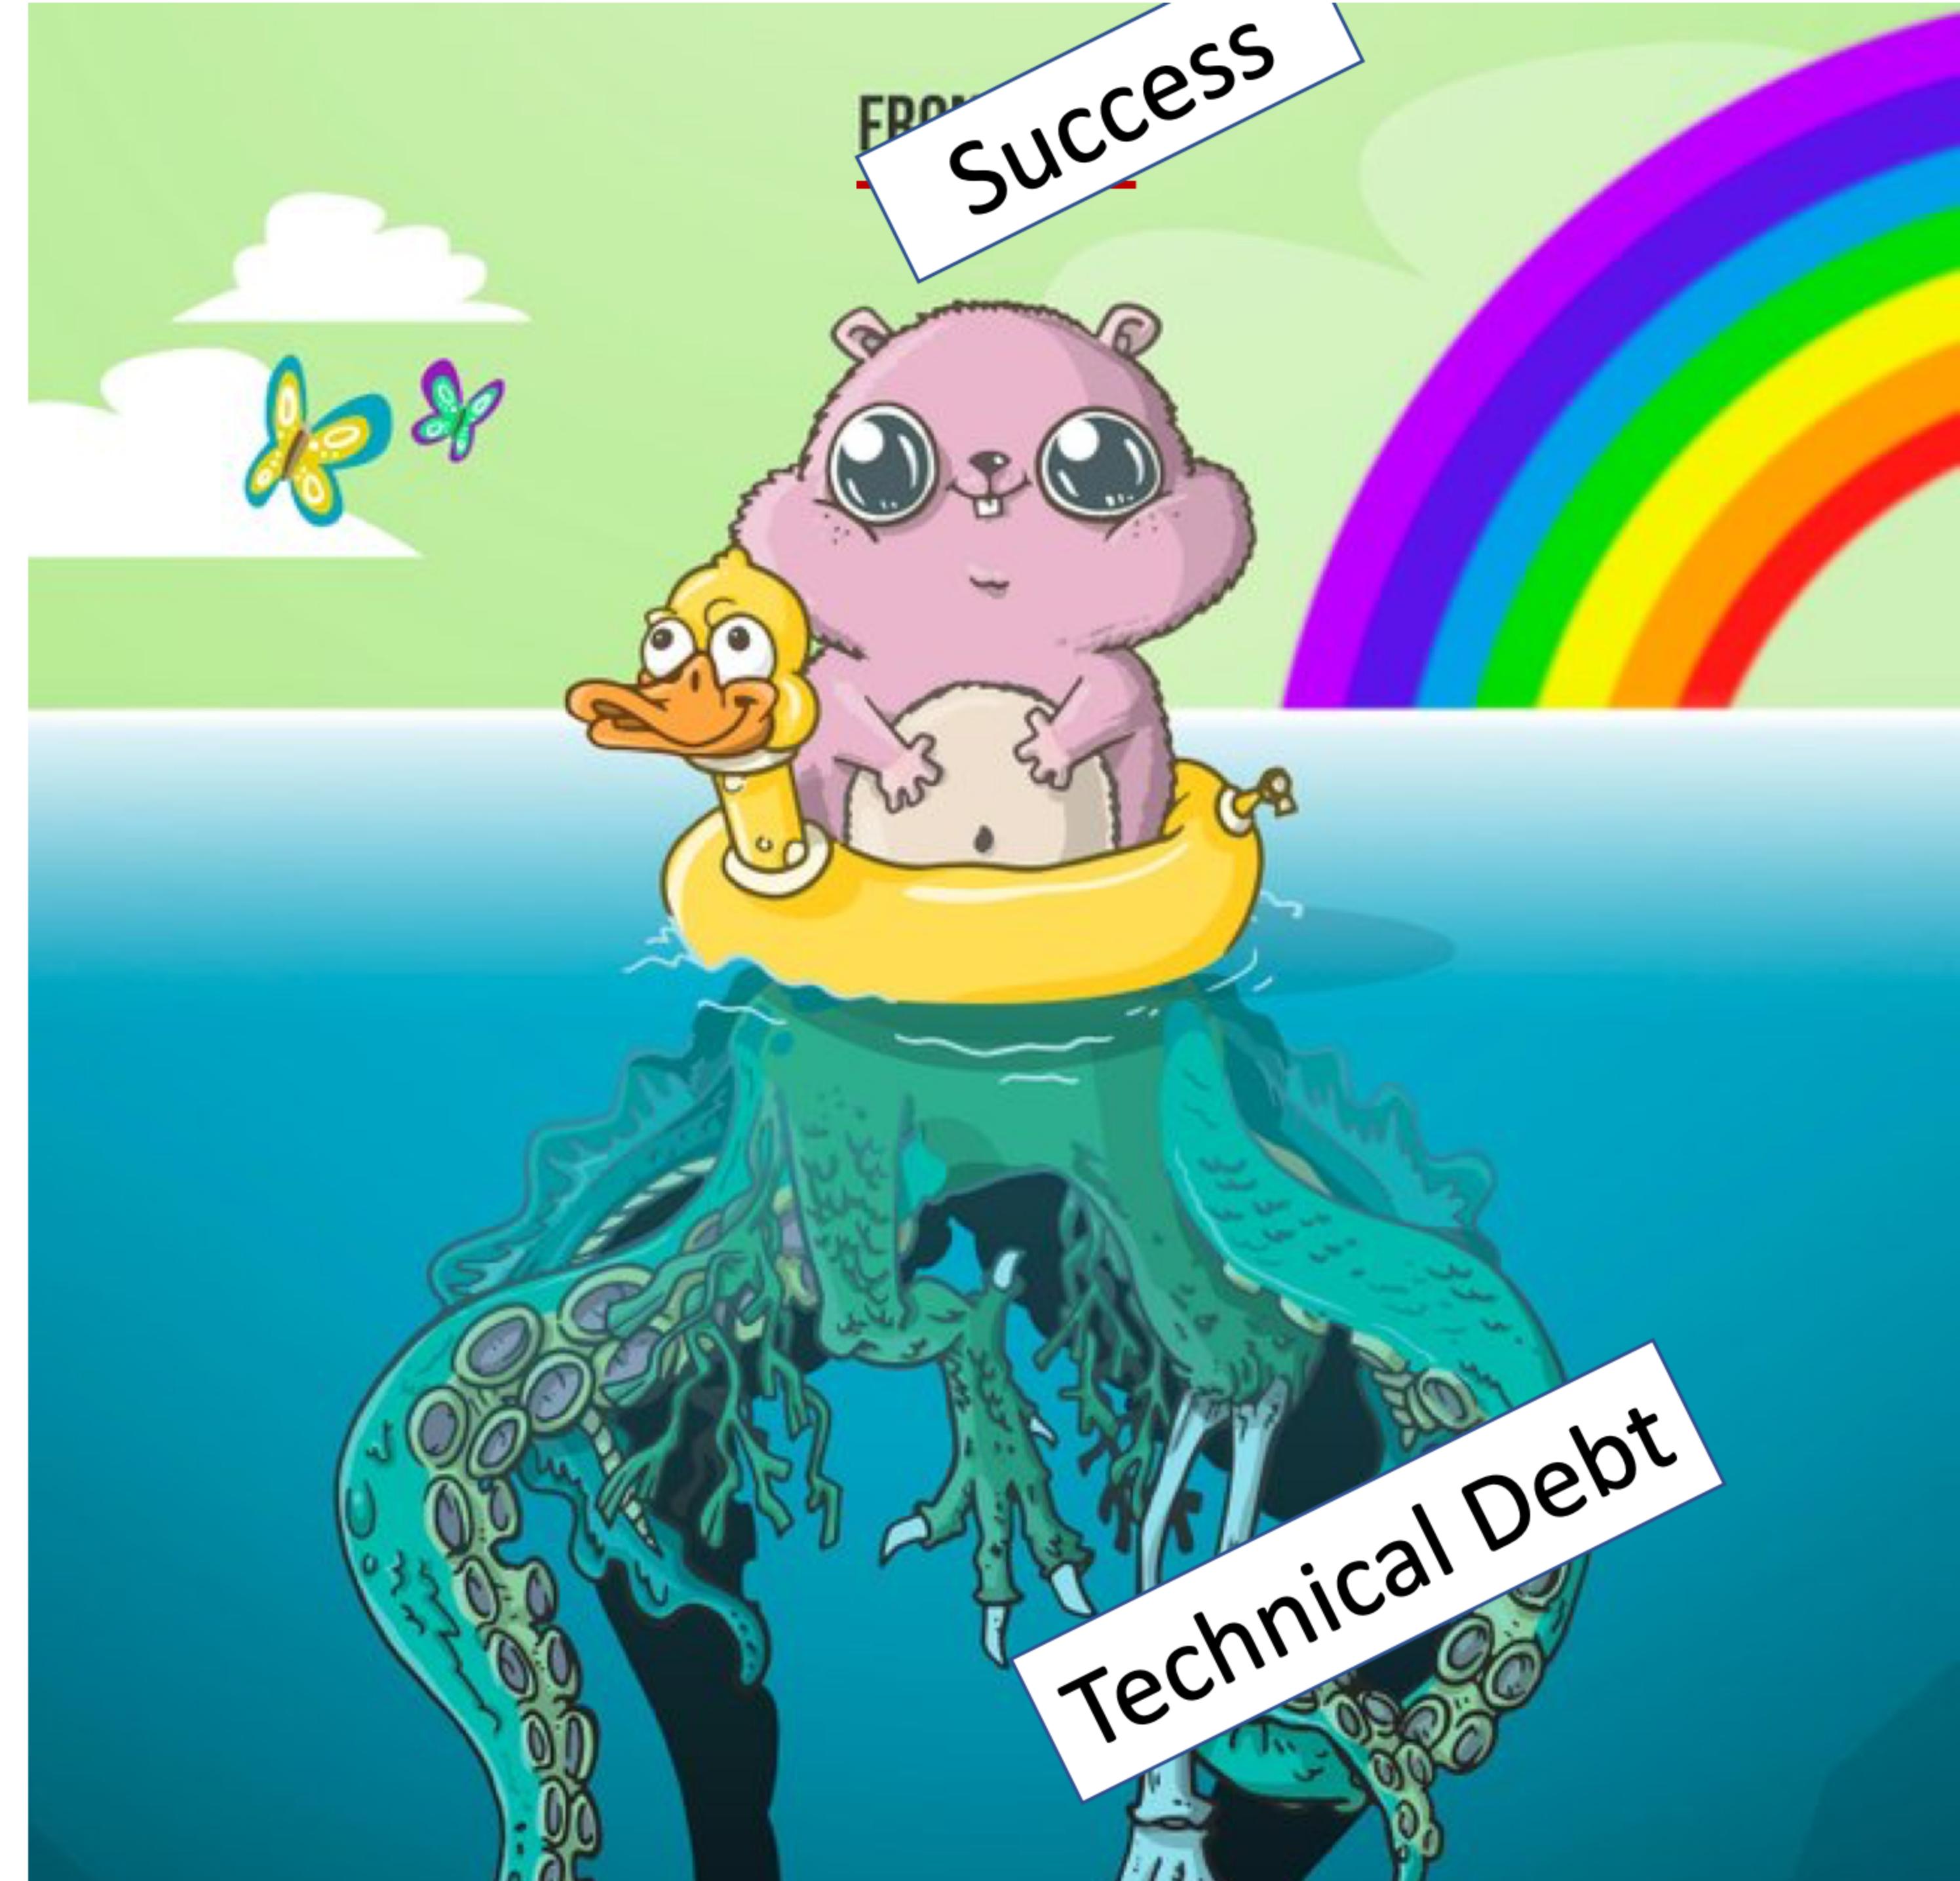
Slow delivery

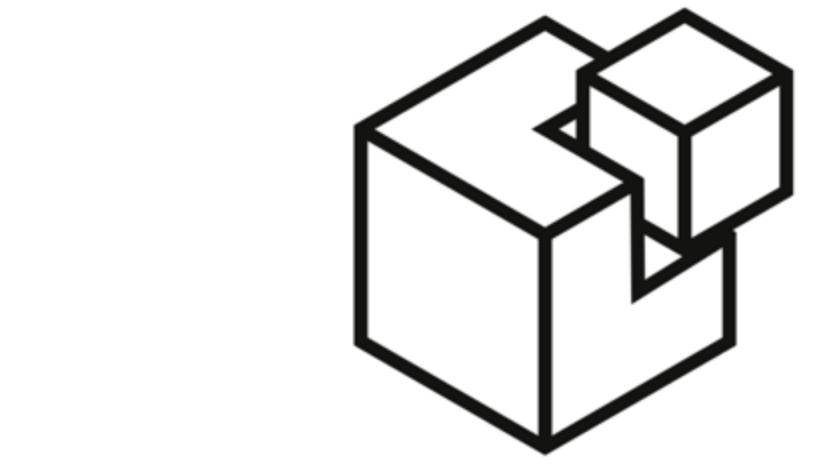
Continuous Deployment?

Fear

Frustration

XX% Idle Machines





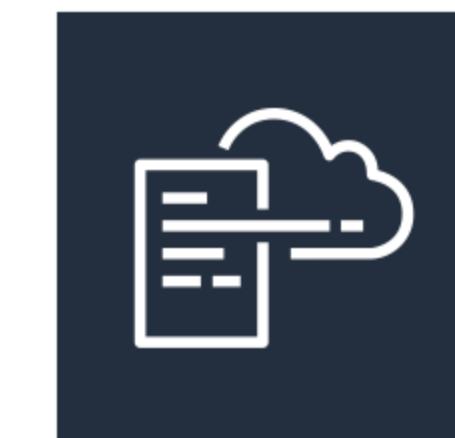
SALTSTACK



ANSIBLE



CHEF



AWS  
CloudFormation

2016



HashiCorp  
**Terraform**

# Black (Blue) Box

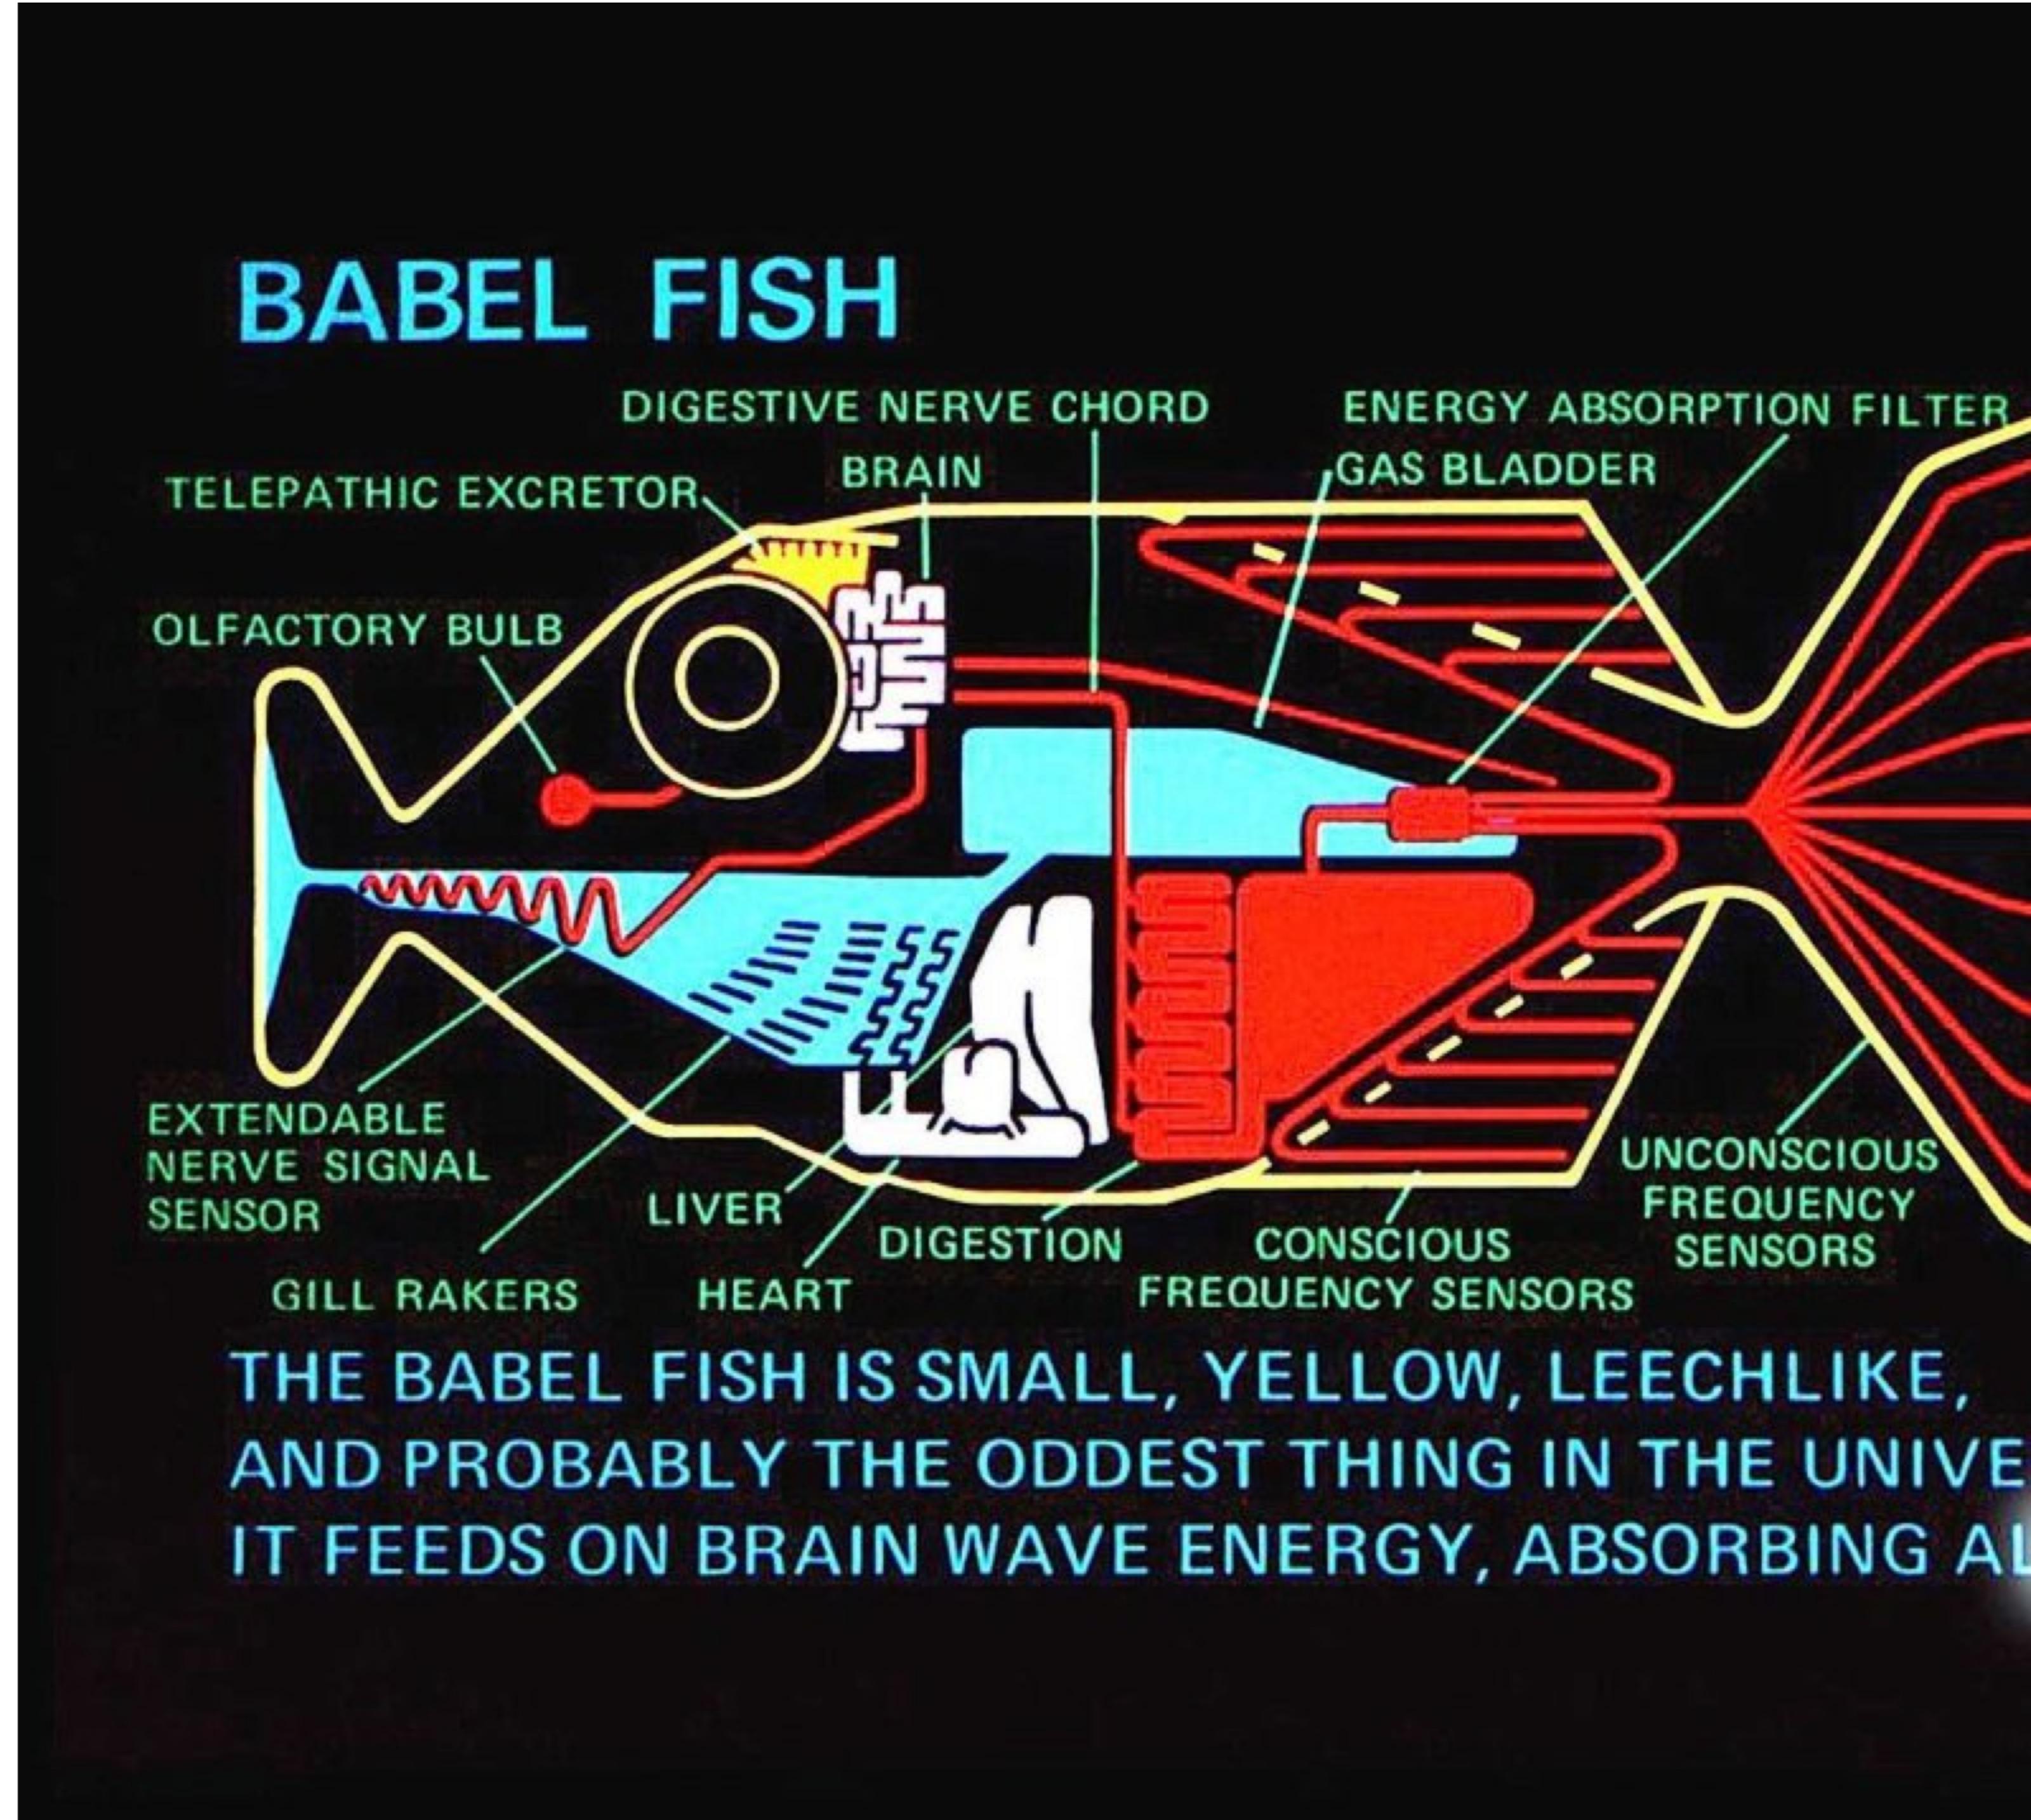
## Infrastructure (almost) invisible

## Easy\* Continuous Deployment



[https://en.wikipedia.org/wiki/File:Dr\\_Who\\_\(316350537\).jpg](https://en.wikipedia.org/wiki/File:Dr_Who_(316350537).jpg)

Common  
Language  
Artifacts  
Platform



# From monolit



# To {micro,macro}services

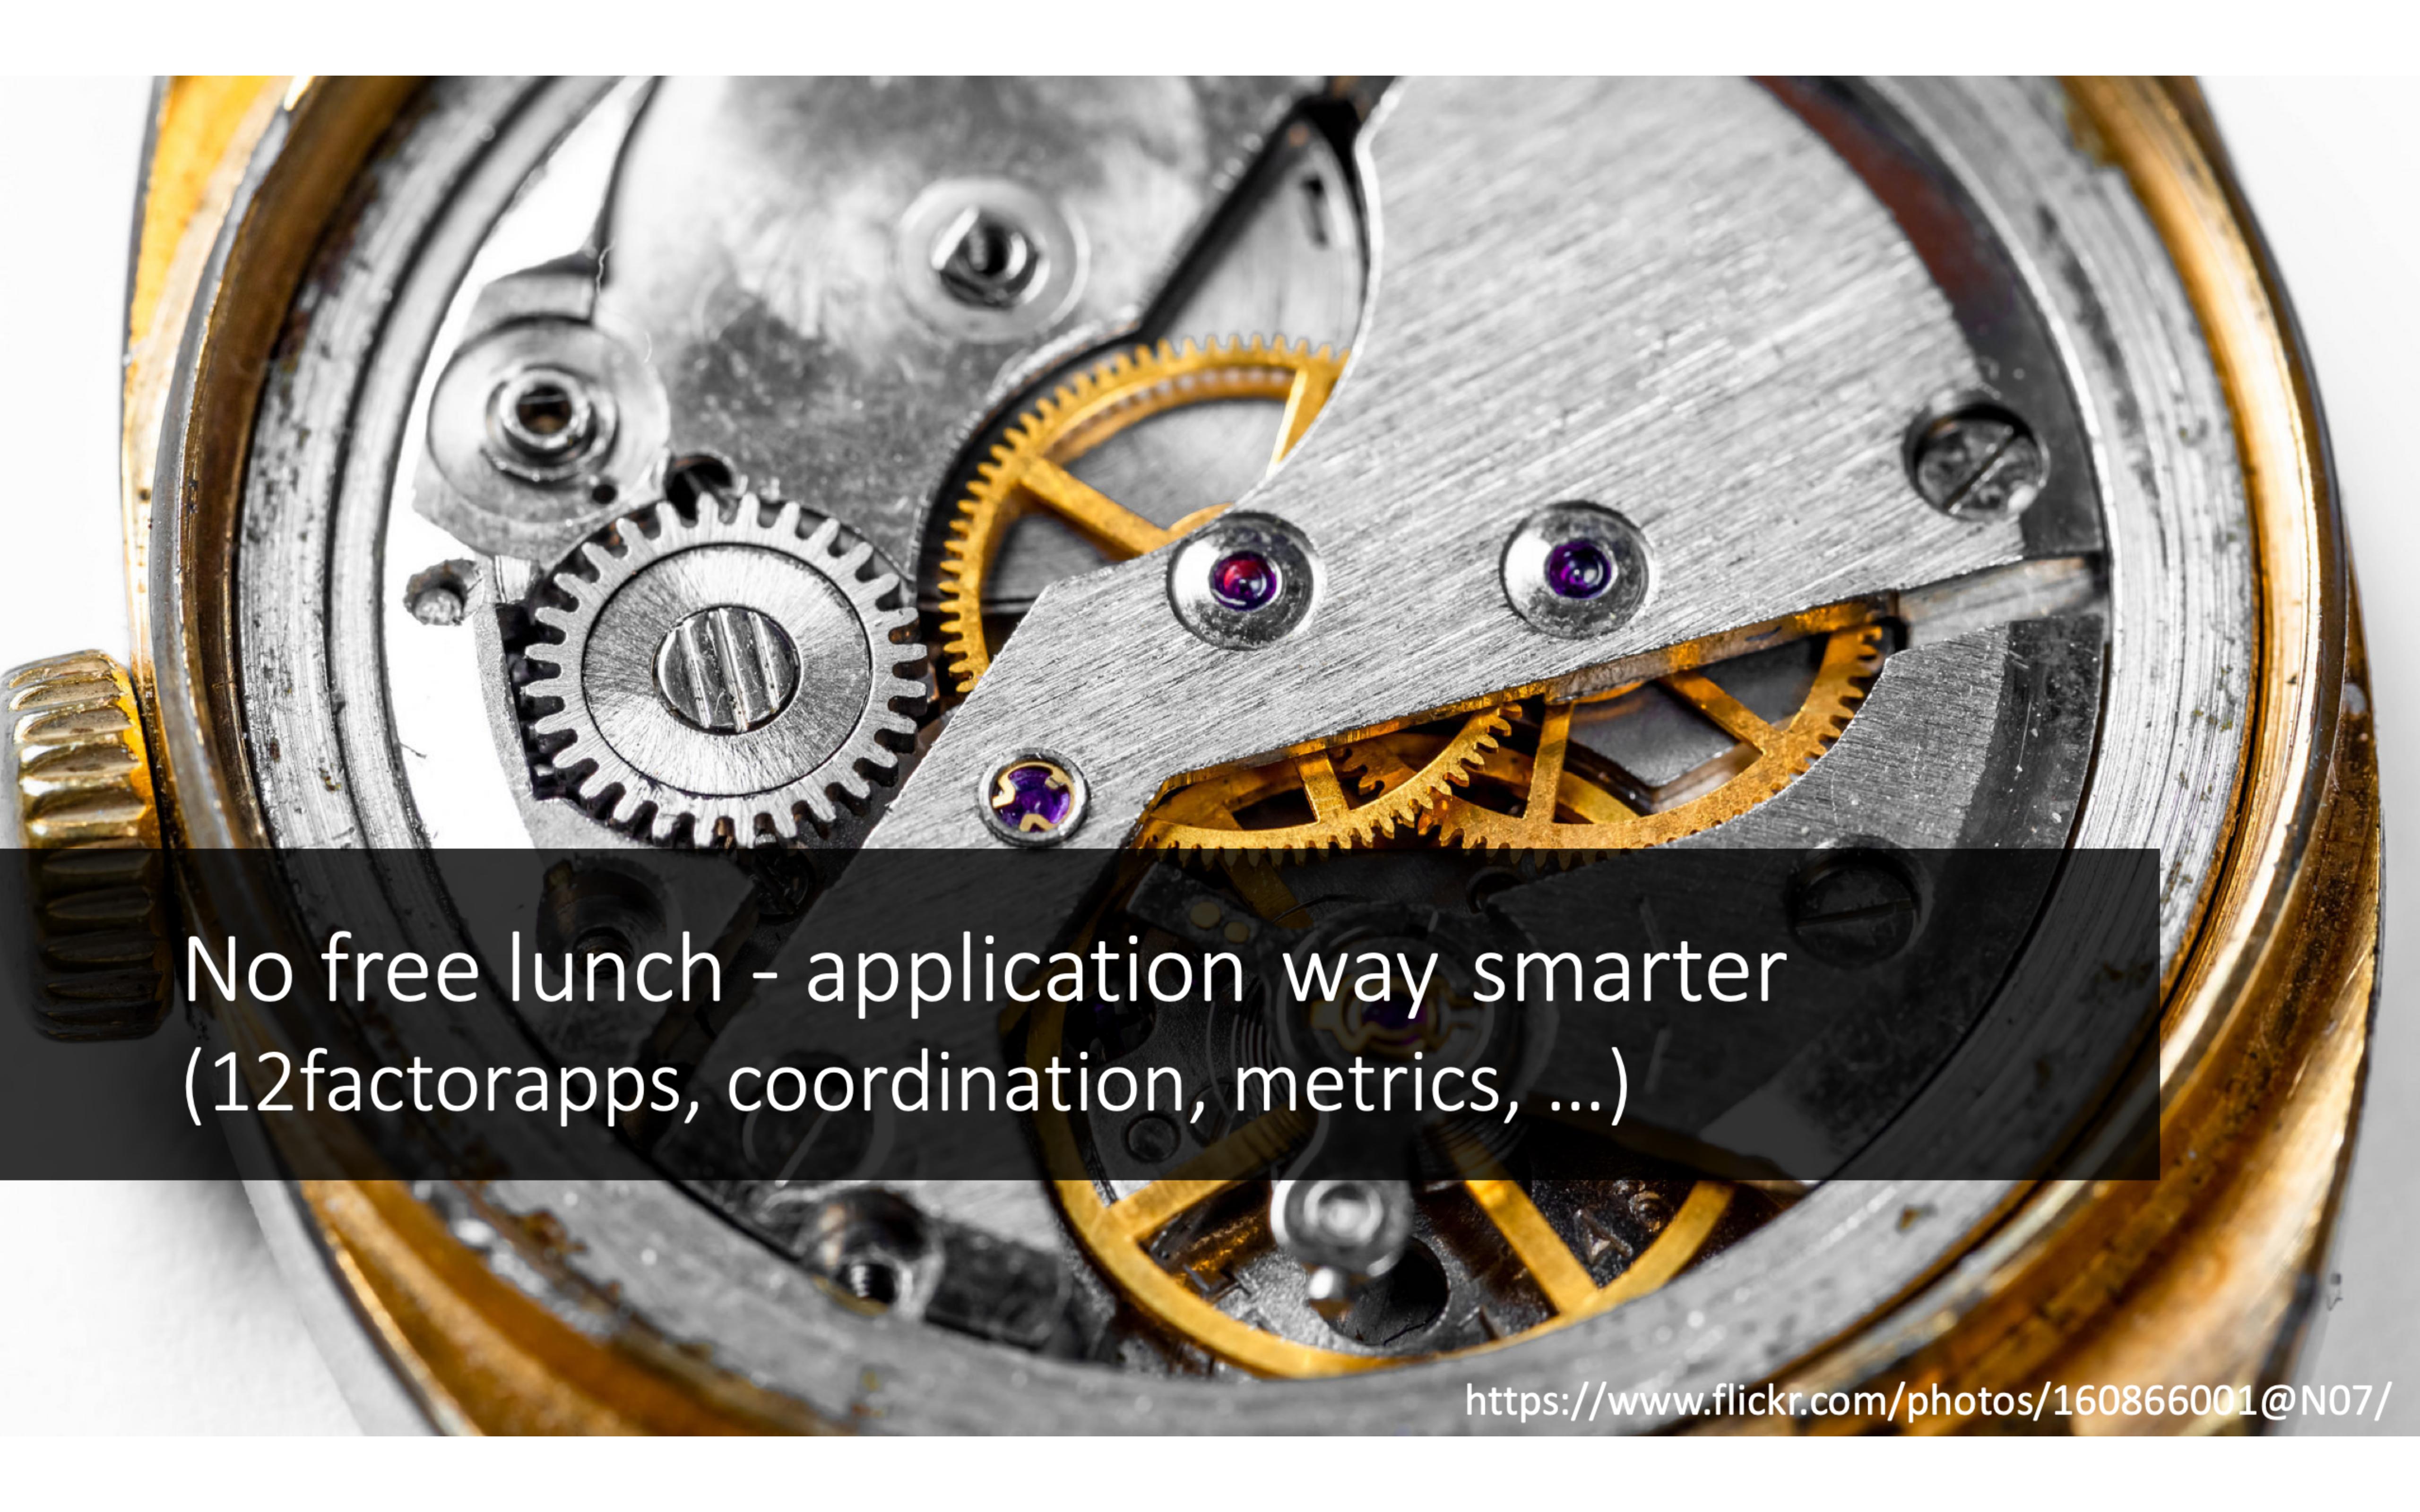


# Learn-as-you-go

1. Deploy Cloud-Native app
2. Make a Hell of Mistakes
3. Get it right or Postpone



envoy



No free lunch - application way smarter  
(12factorapps, coordination, metrics, ...)

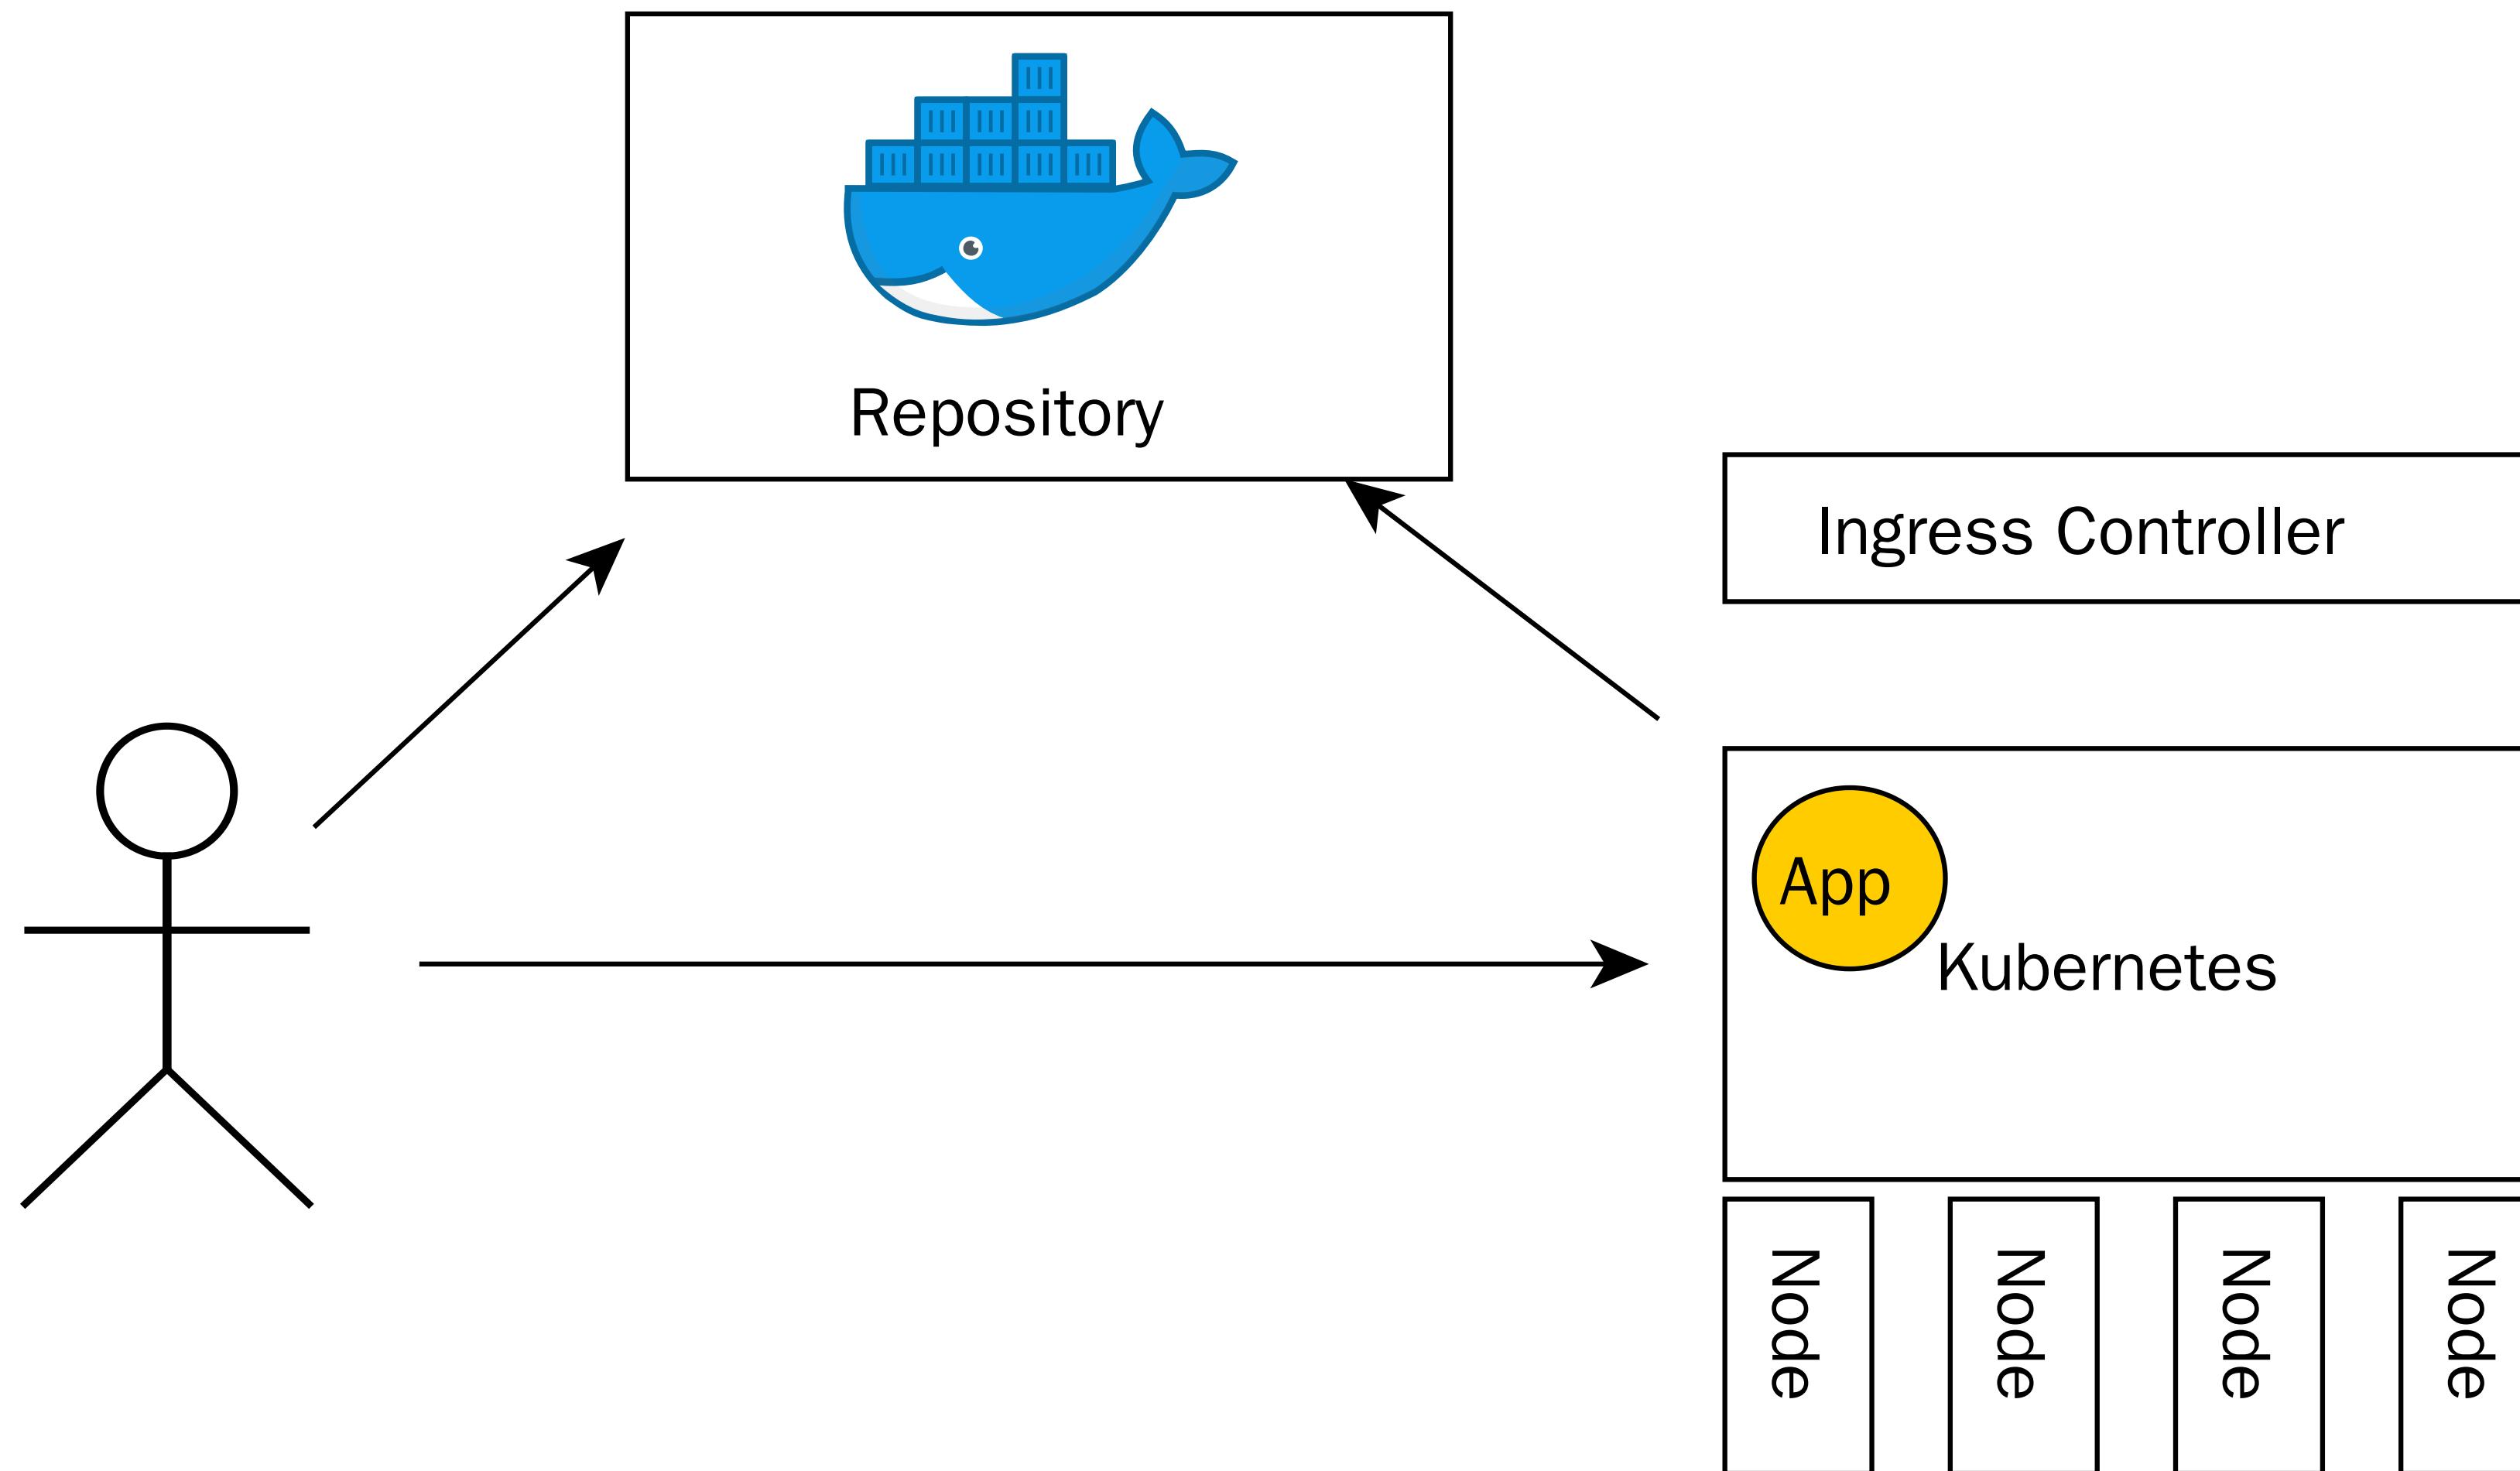
# Kubernetes

- Container management
- Battery for 12factor apps
- Decouple dev and admin work

## Battery for 12factor apps

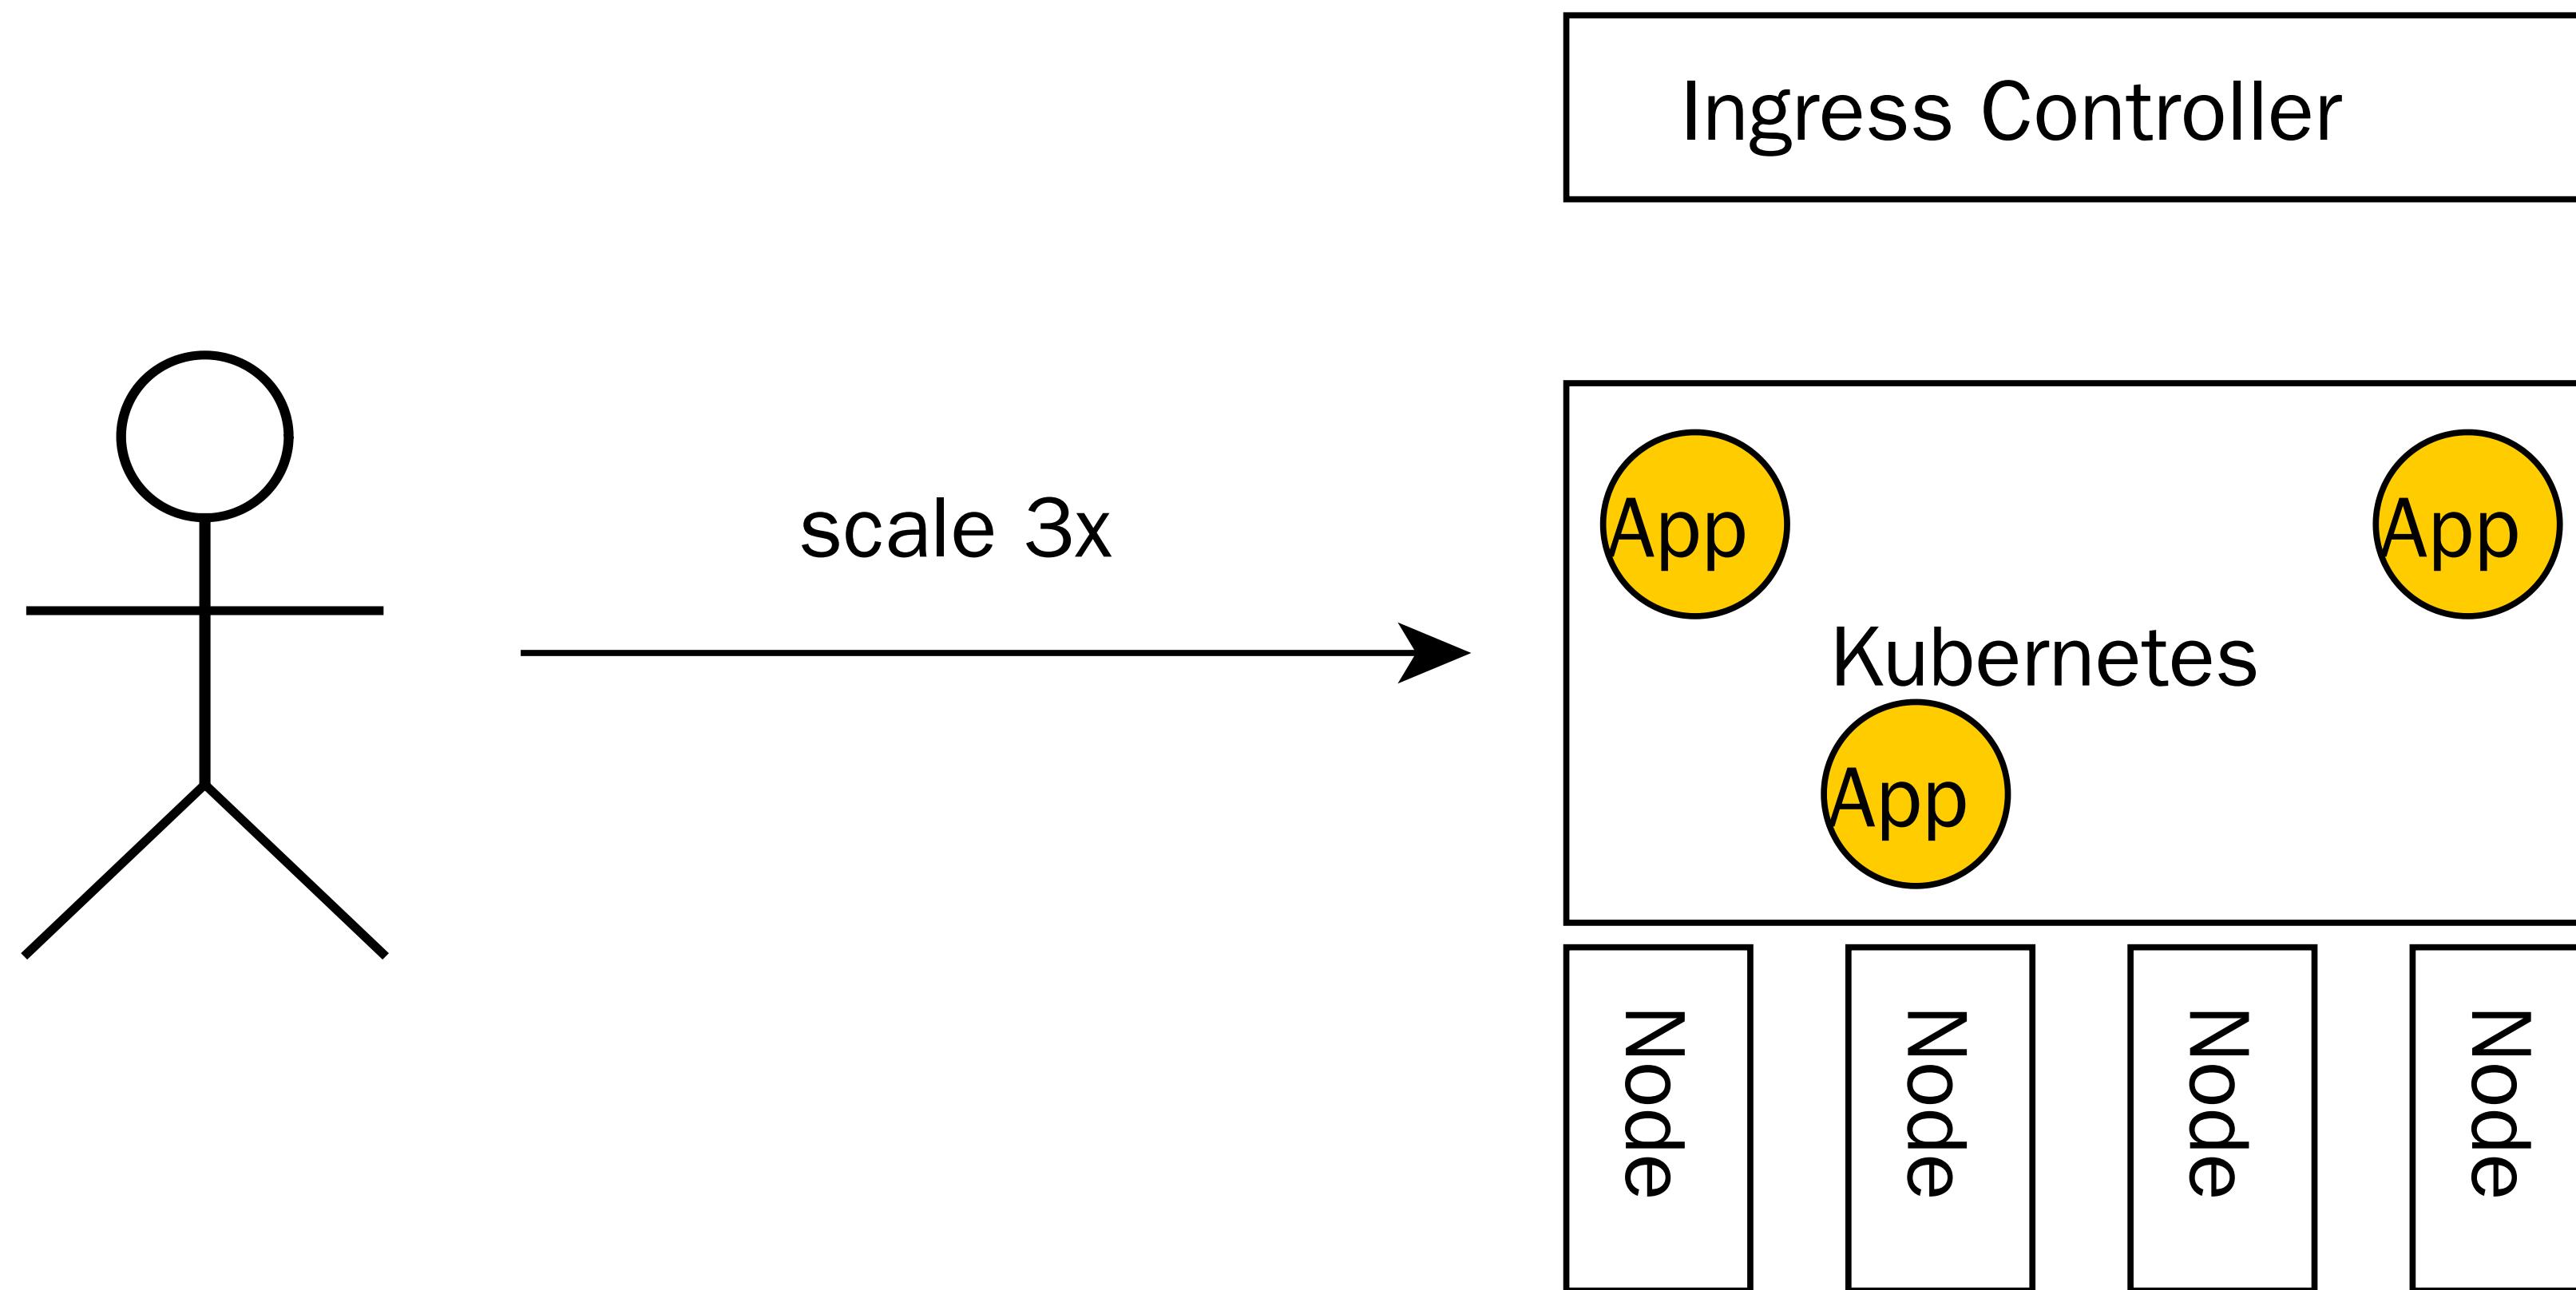
- Service Discovery
- Load Balancing
- Automatic restarts
- Lingua Franca

# KUBERNETES



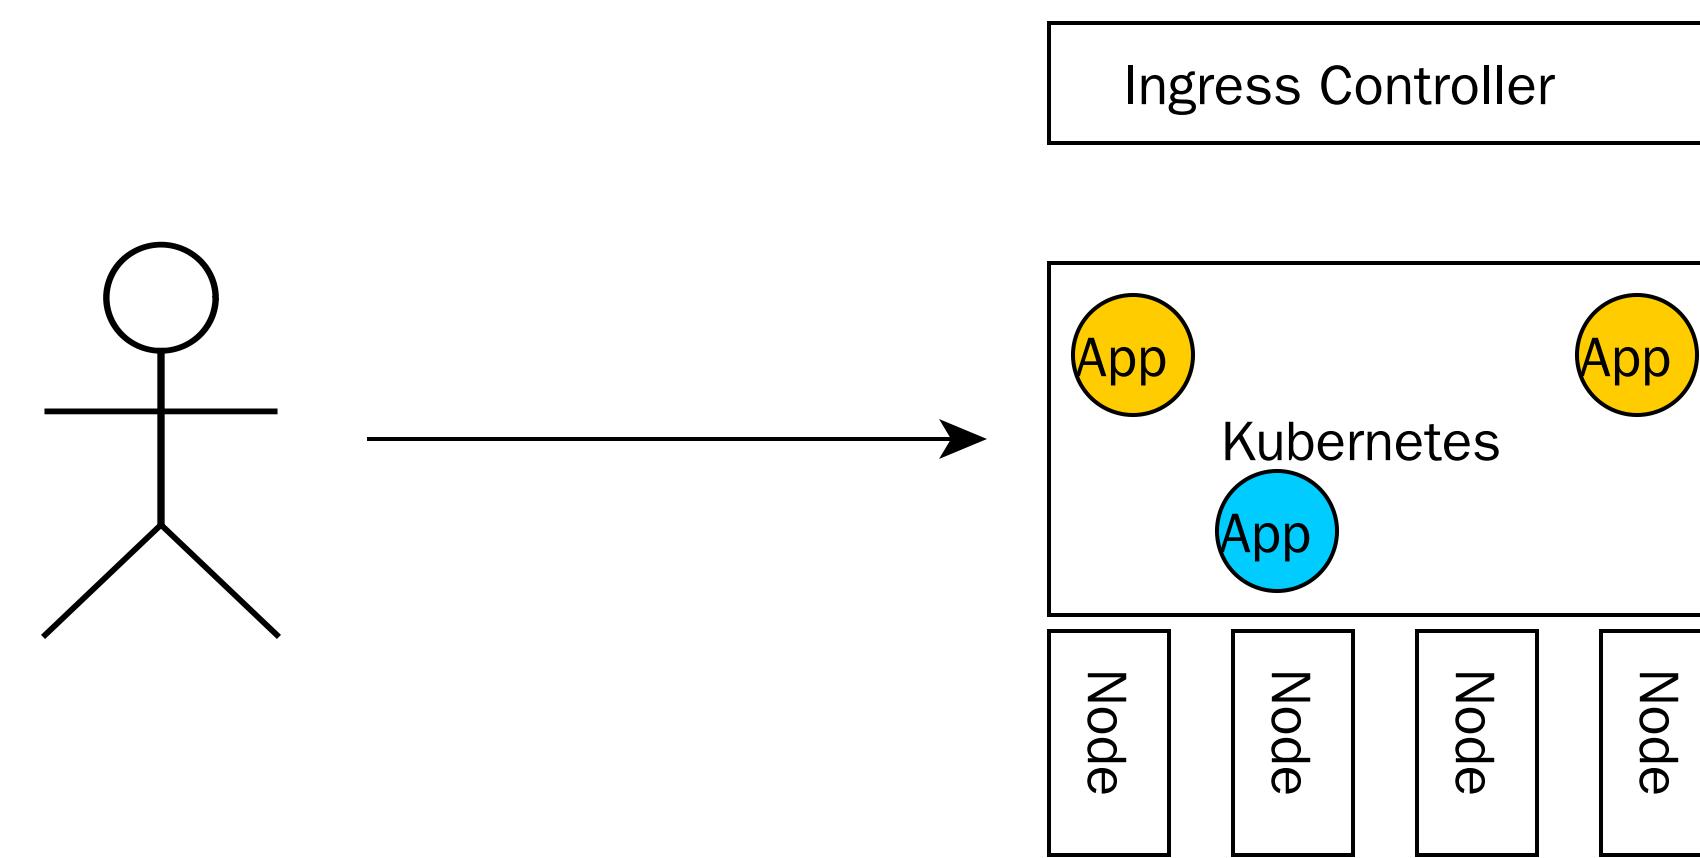
make docker\_push; kubectl create -f app-srv-dpl.yaml

# Scale up! Scale down!

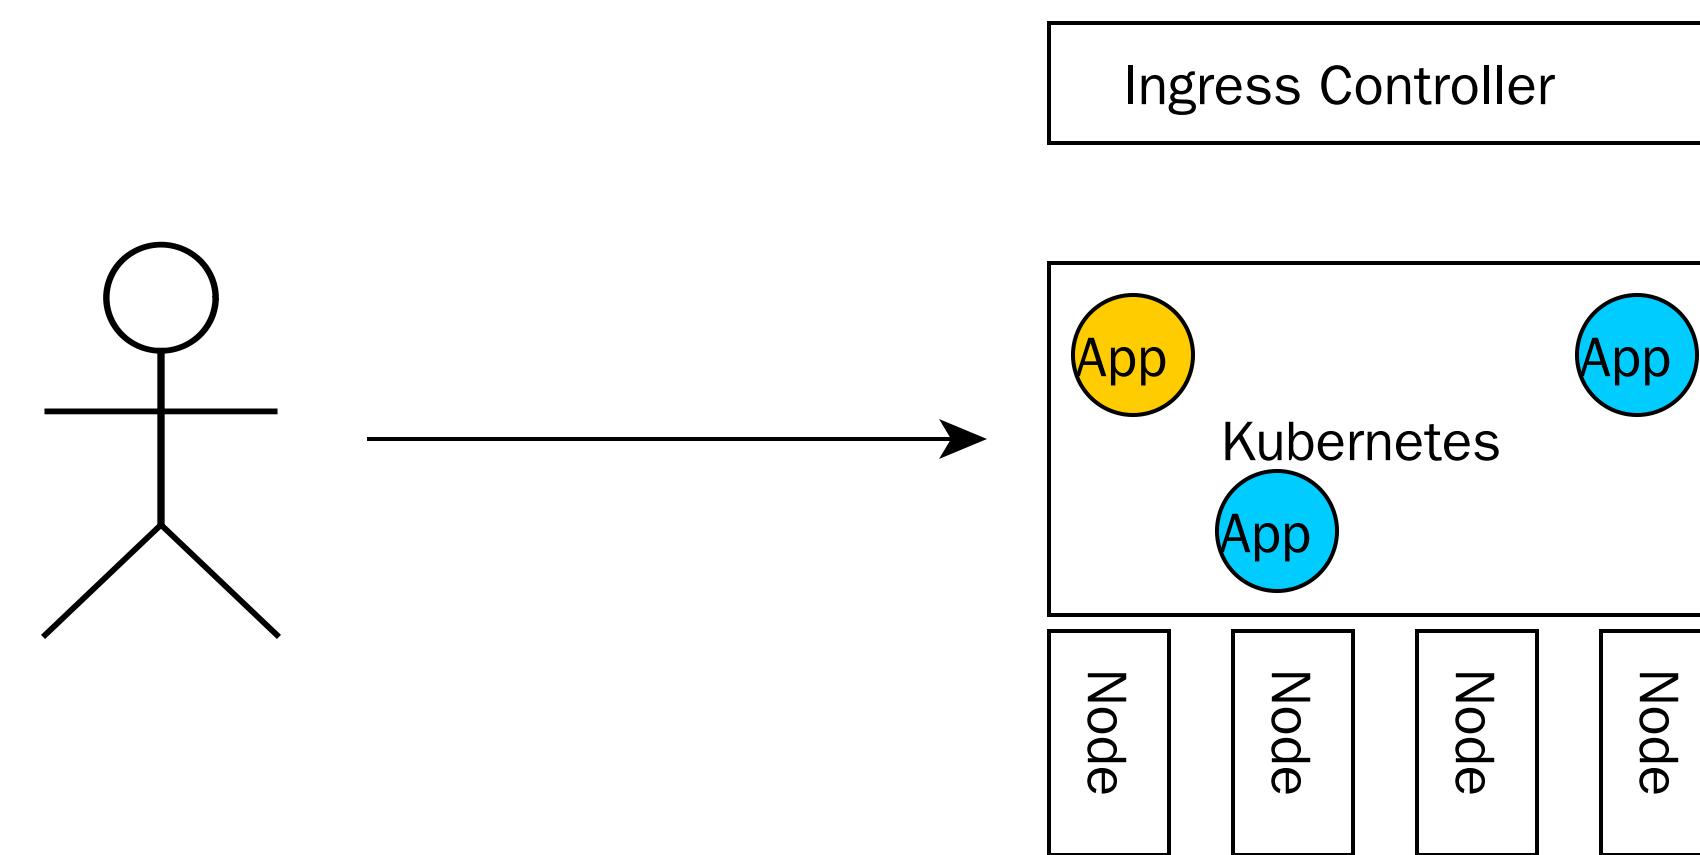


```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

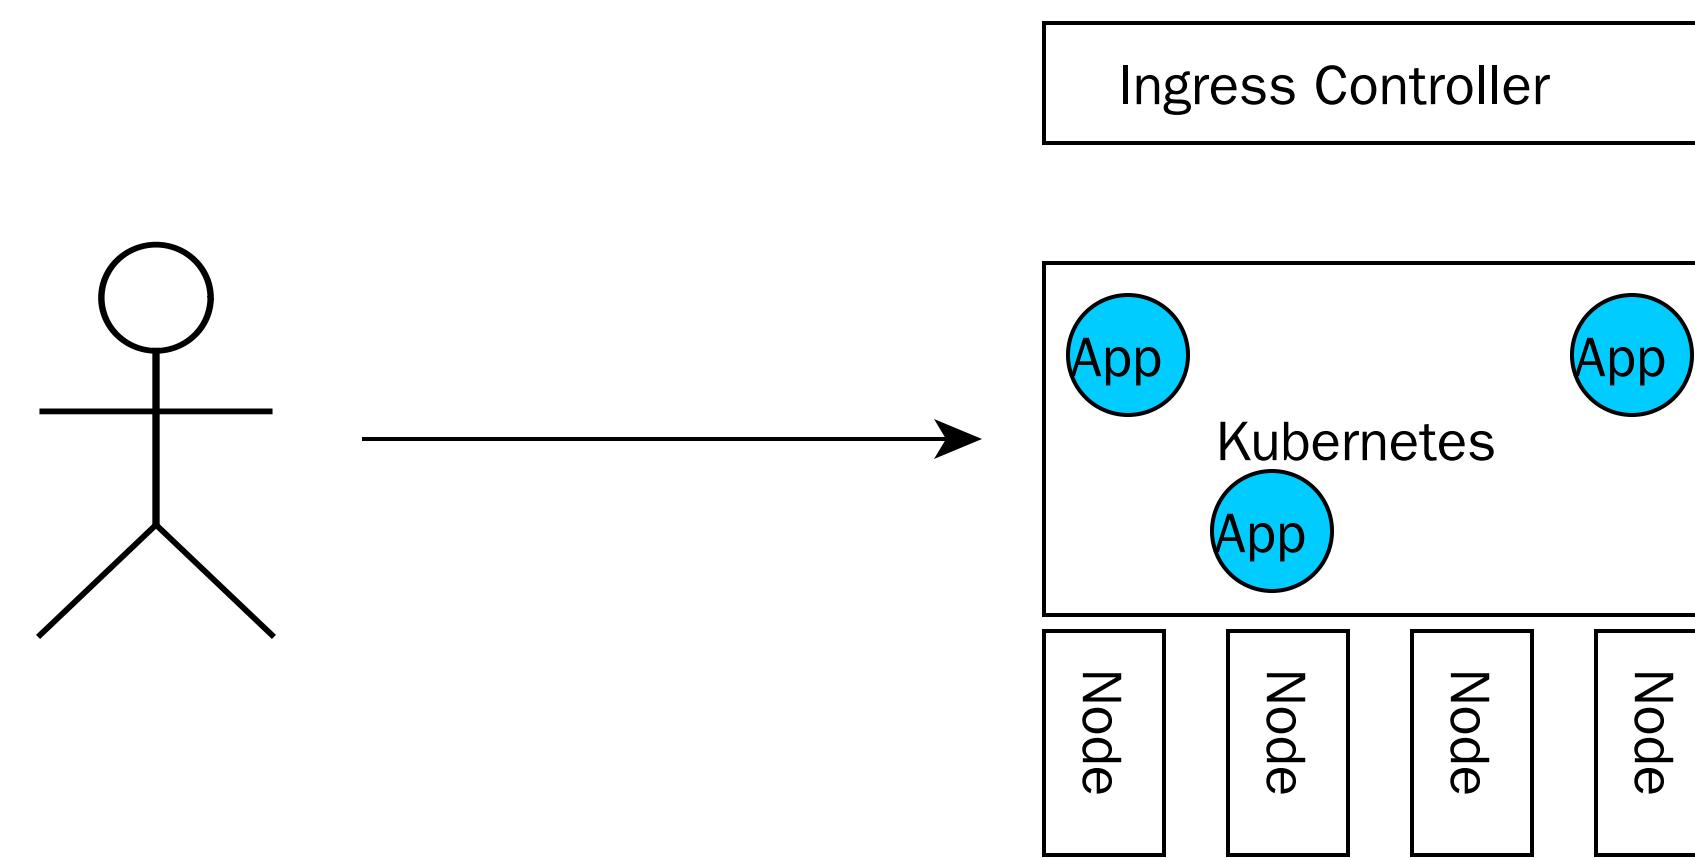
# Example: rolling updates



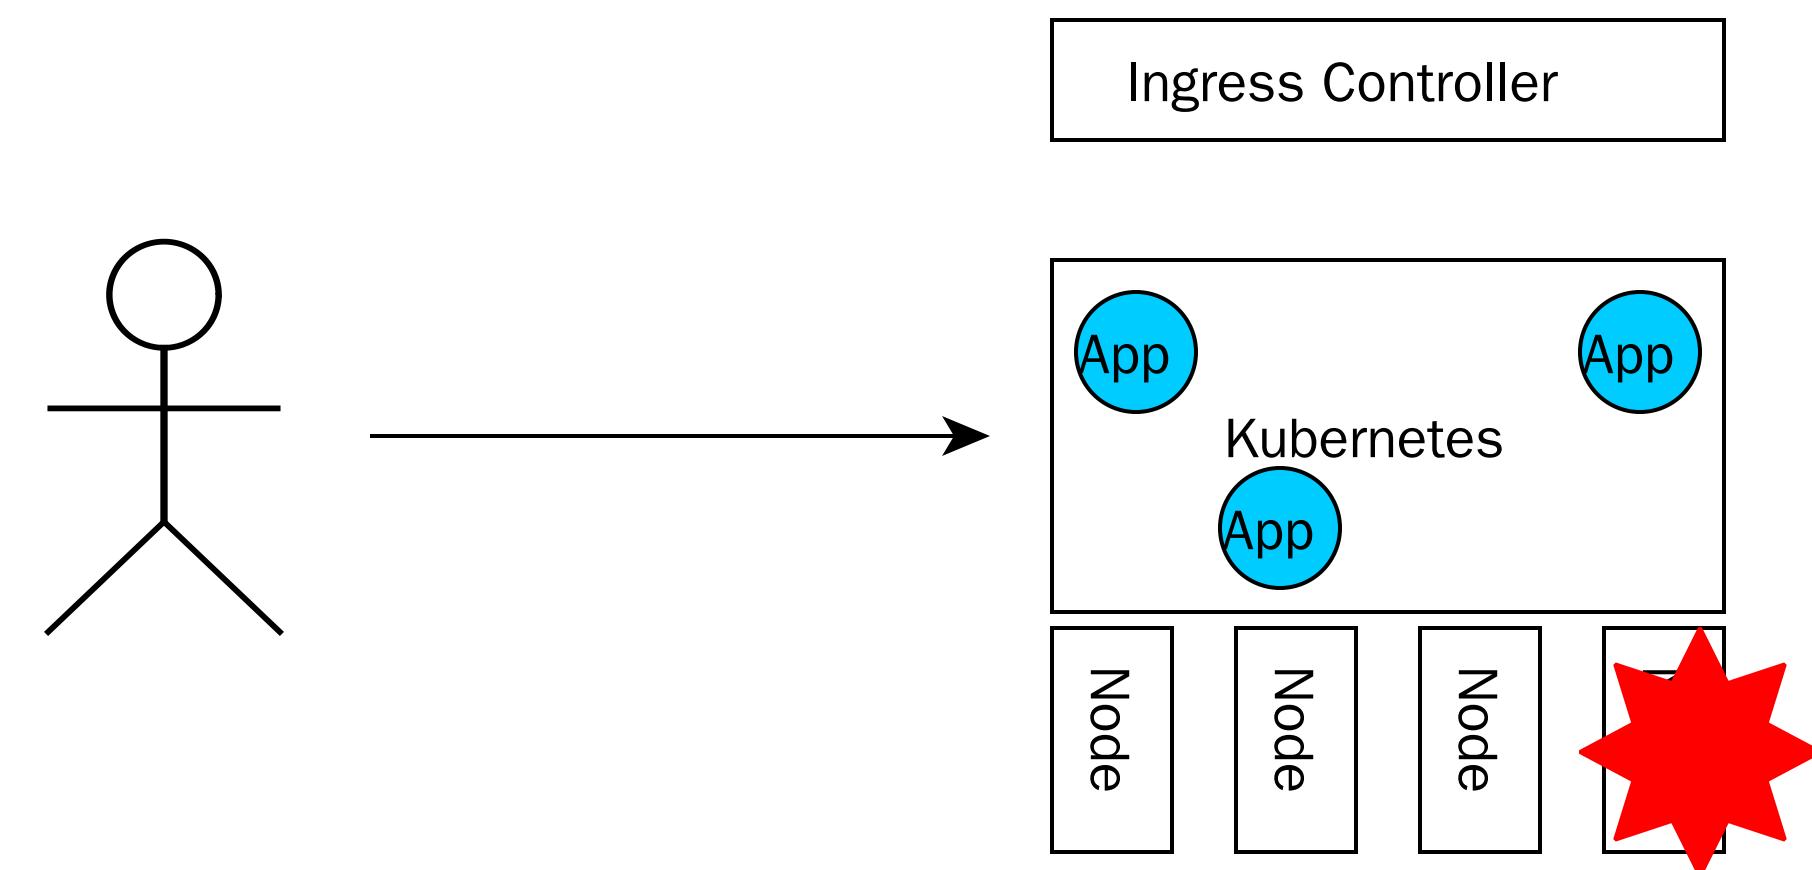
# Example: rolling updates



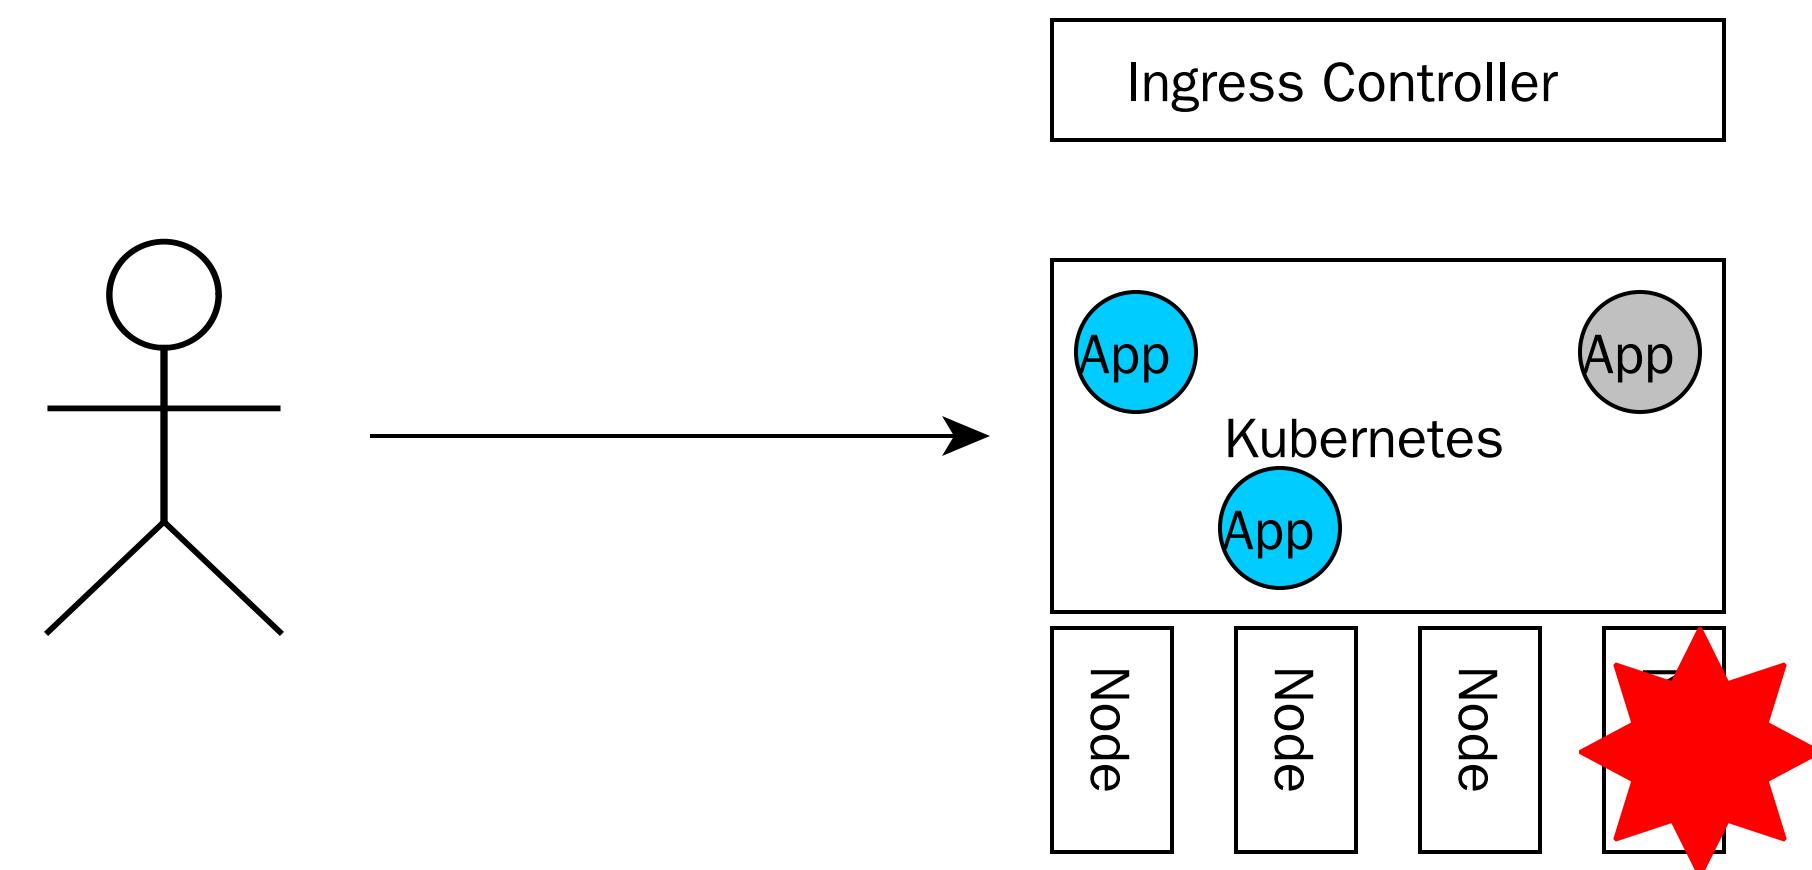
# Example: rolling updates



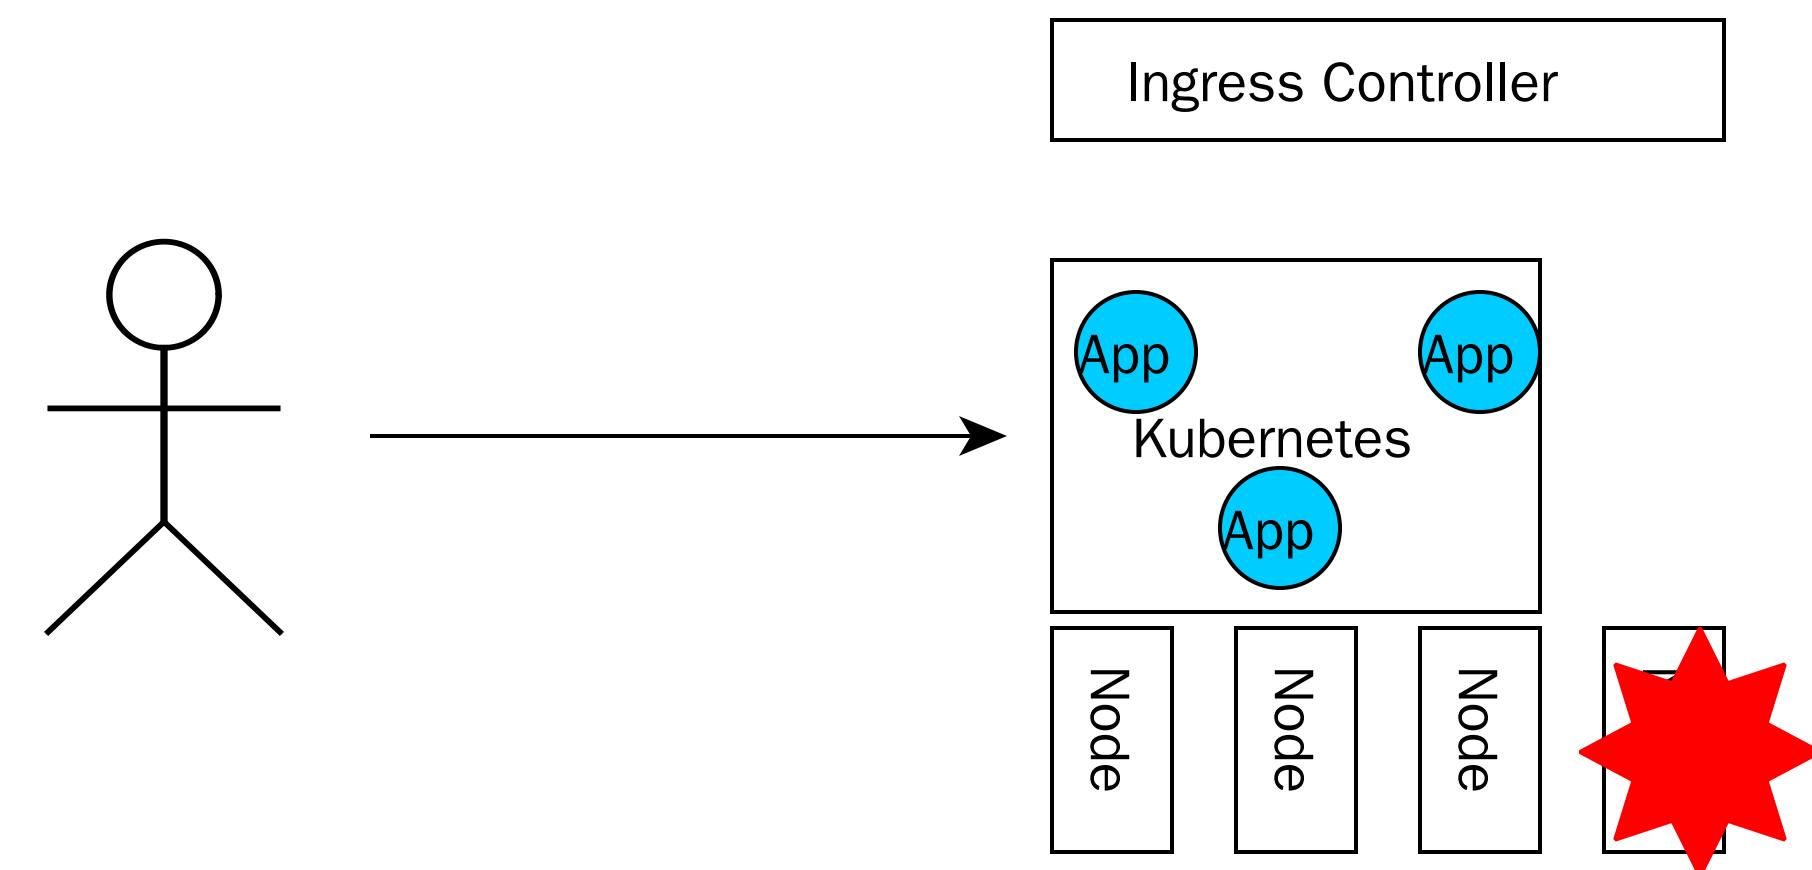
# Resilient infra



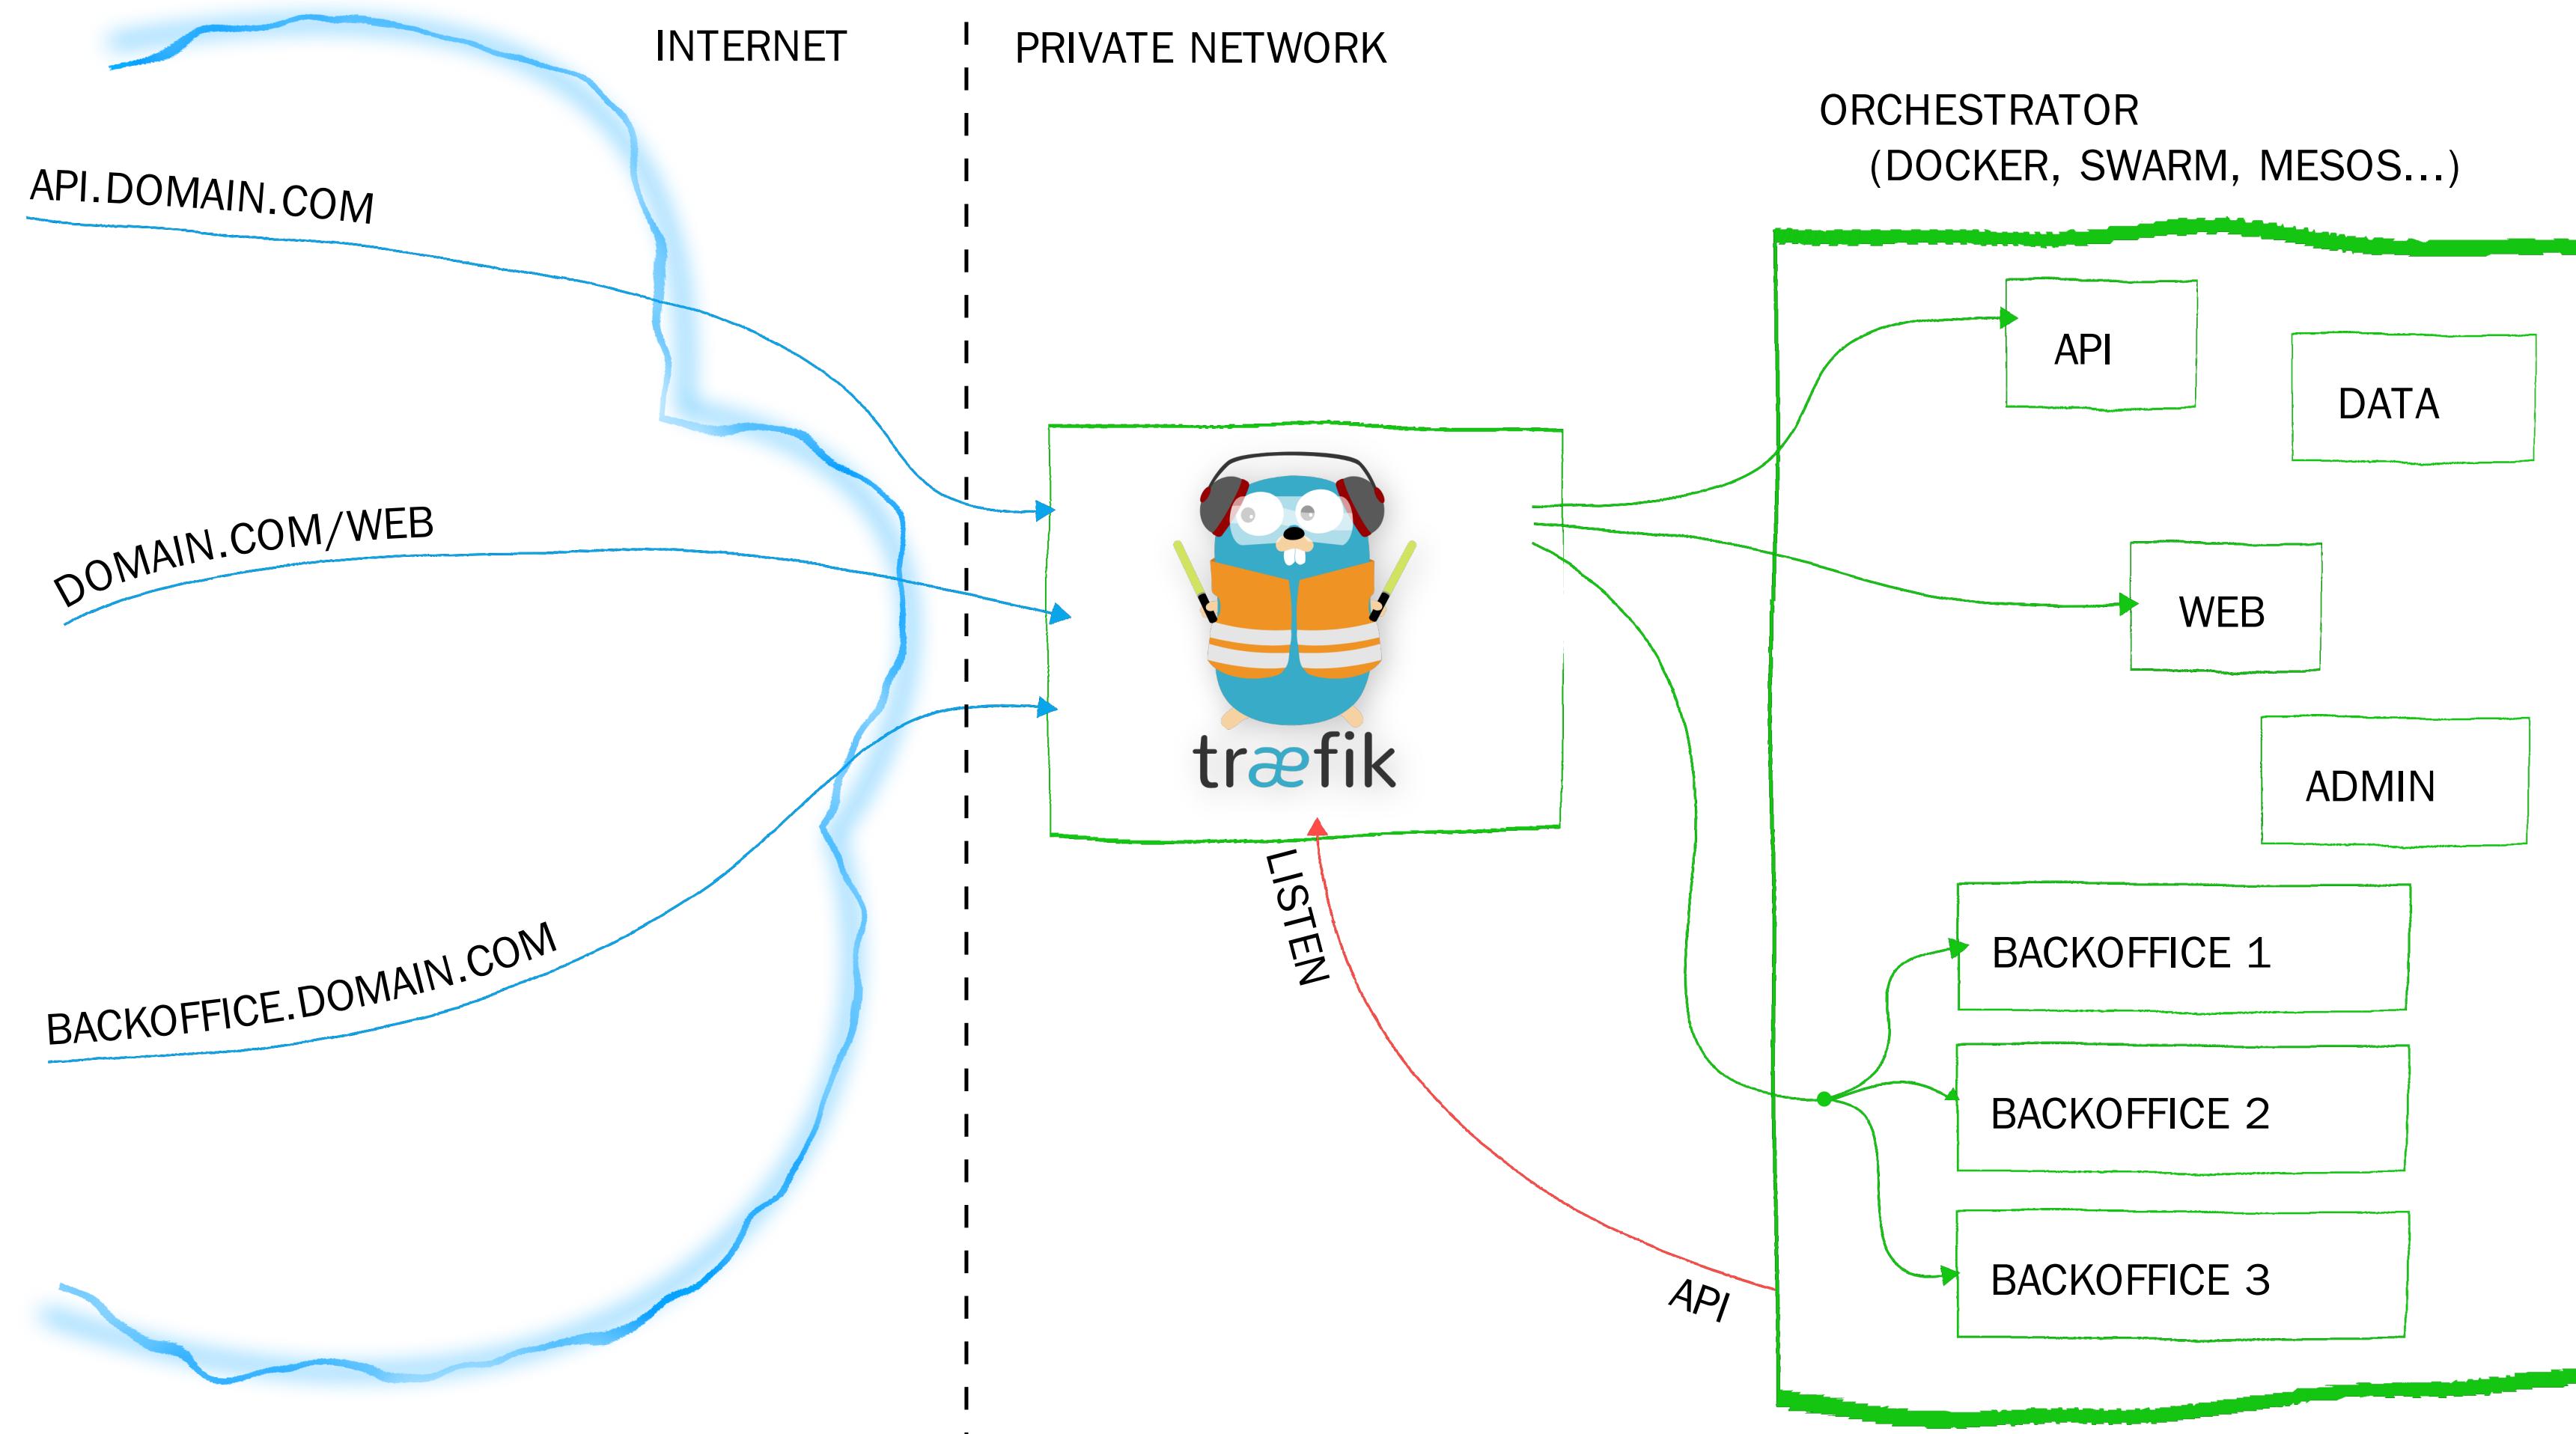
# Resilient infra



# Resilient infra



# Ingress



Pluggable ingress controller.

# Ingress

## Integration with Cloud providers

- LoadBalancer will be provisioned from K8S

# Ingress

## Pattern

api.example.io/v1/users    users-v1

---

api.example.io/v2/users    users-v2

---

## example.io

web

Patterns/Rules support wide routing options.

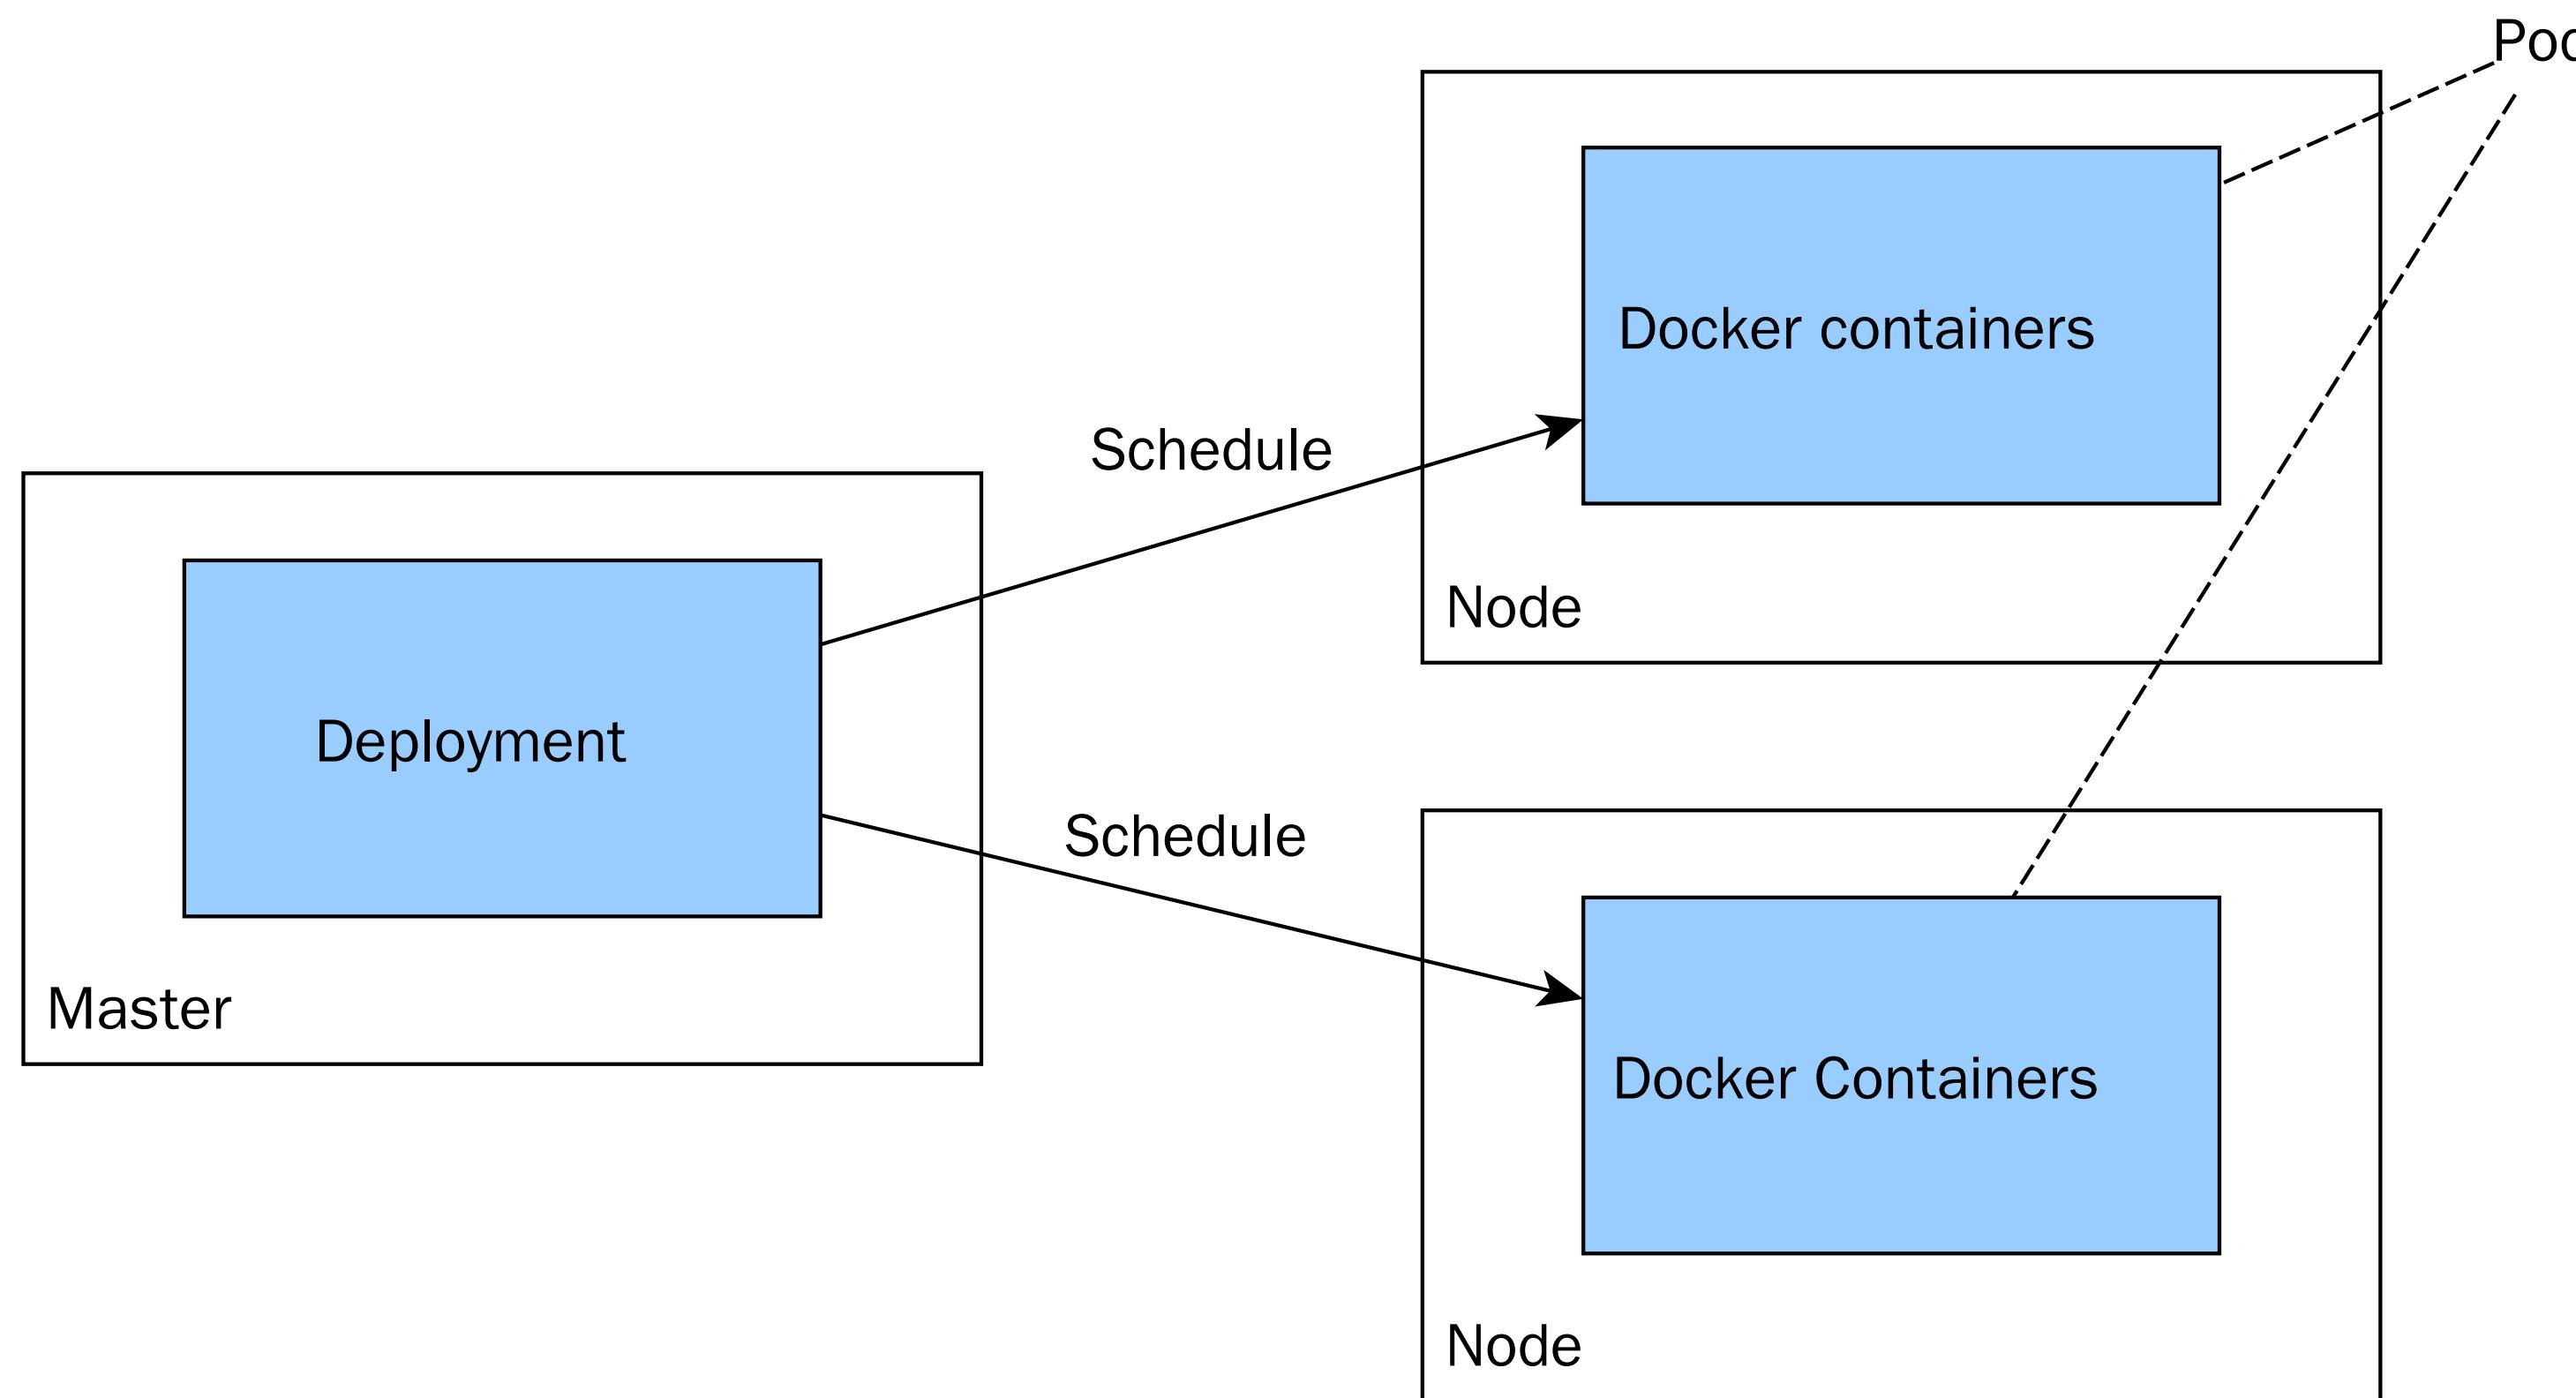
## Basic concepts

Name	Purpose	
Ingress	Customer interface	URL to Service

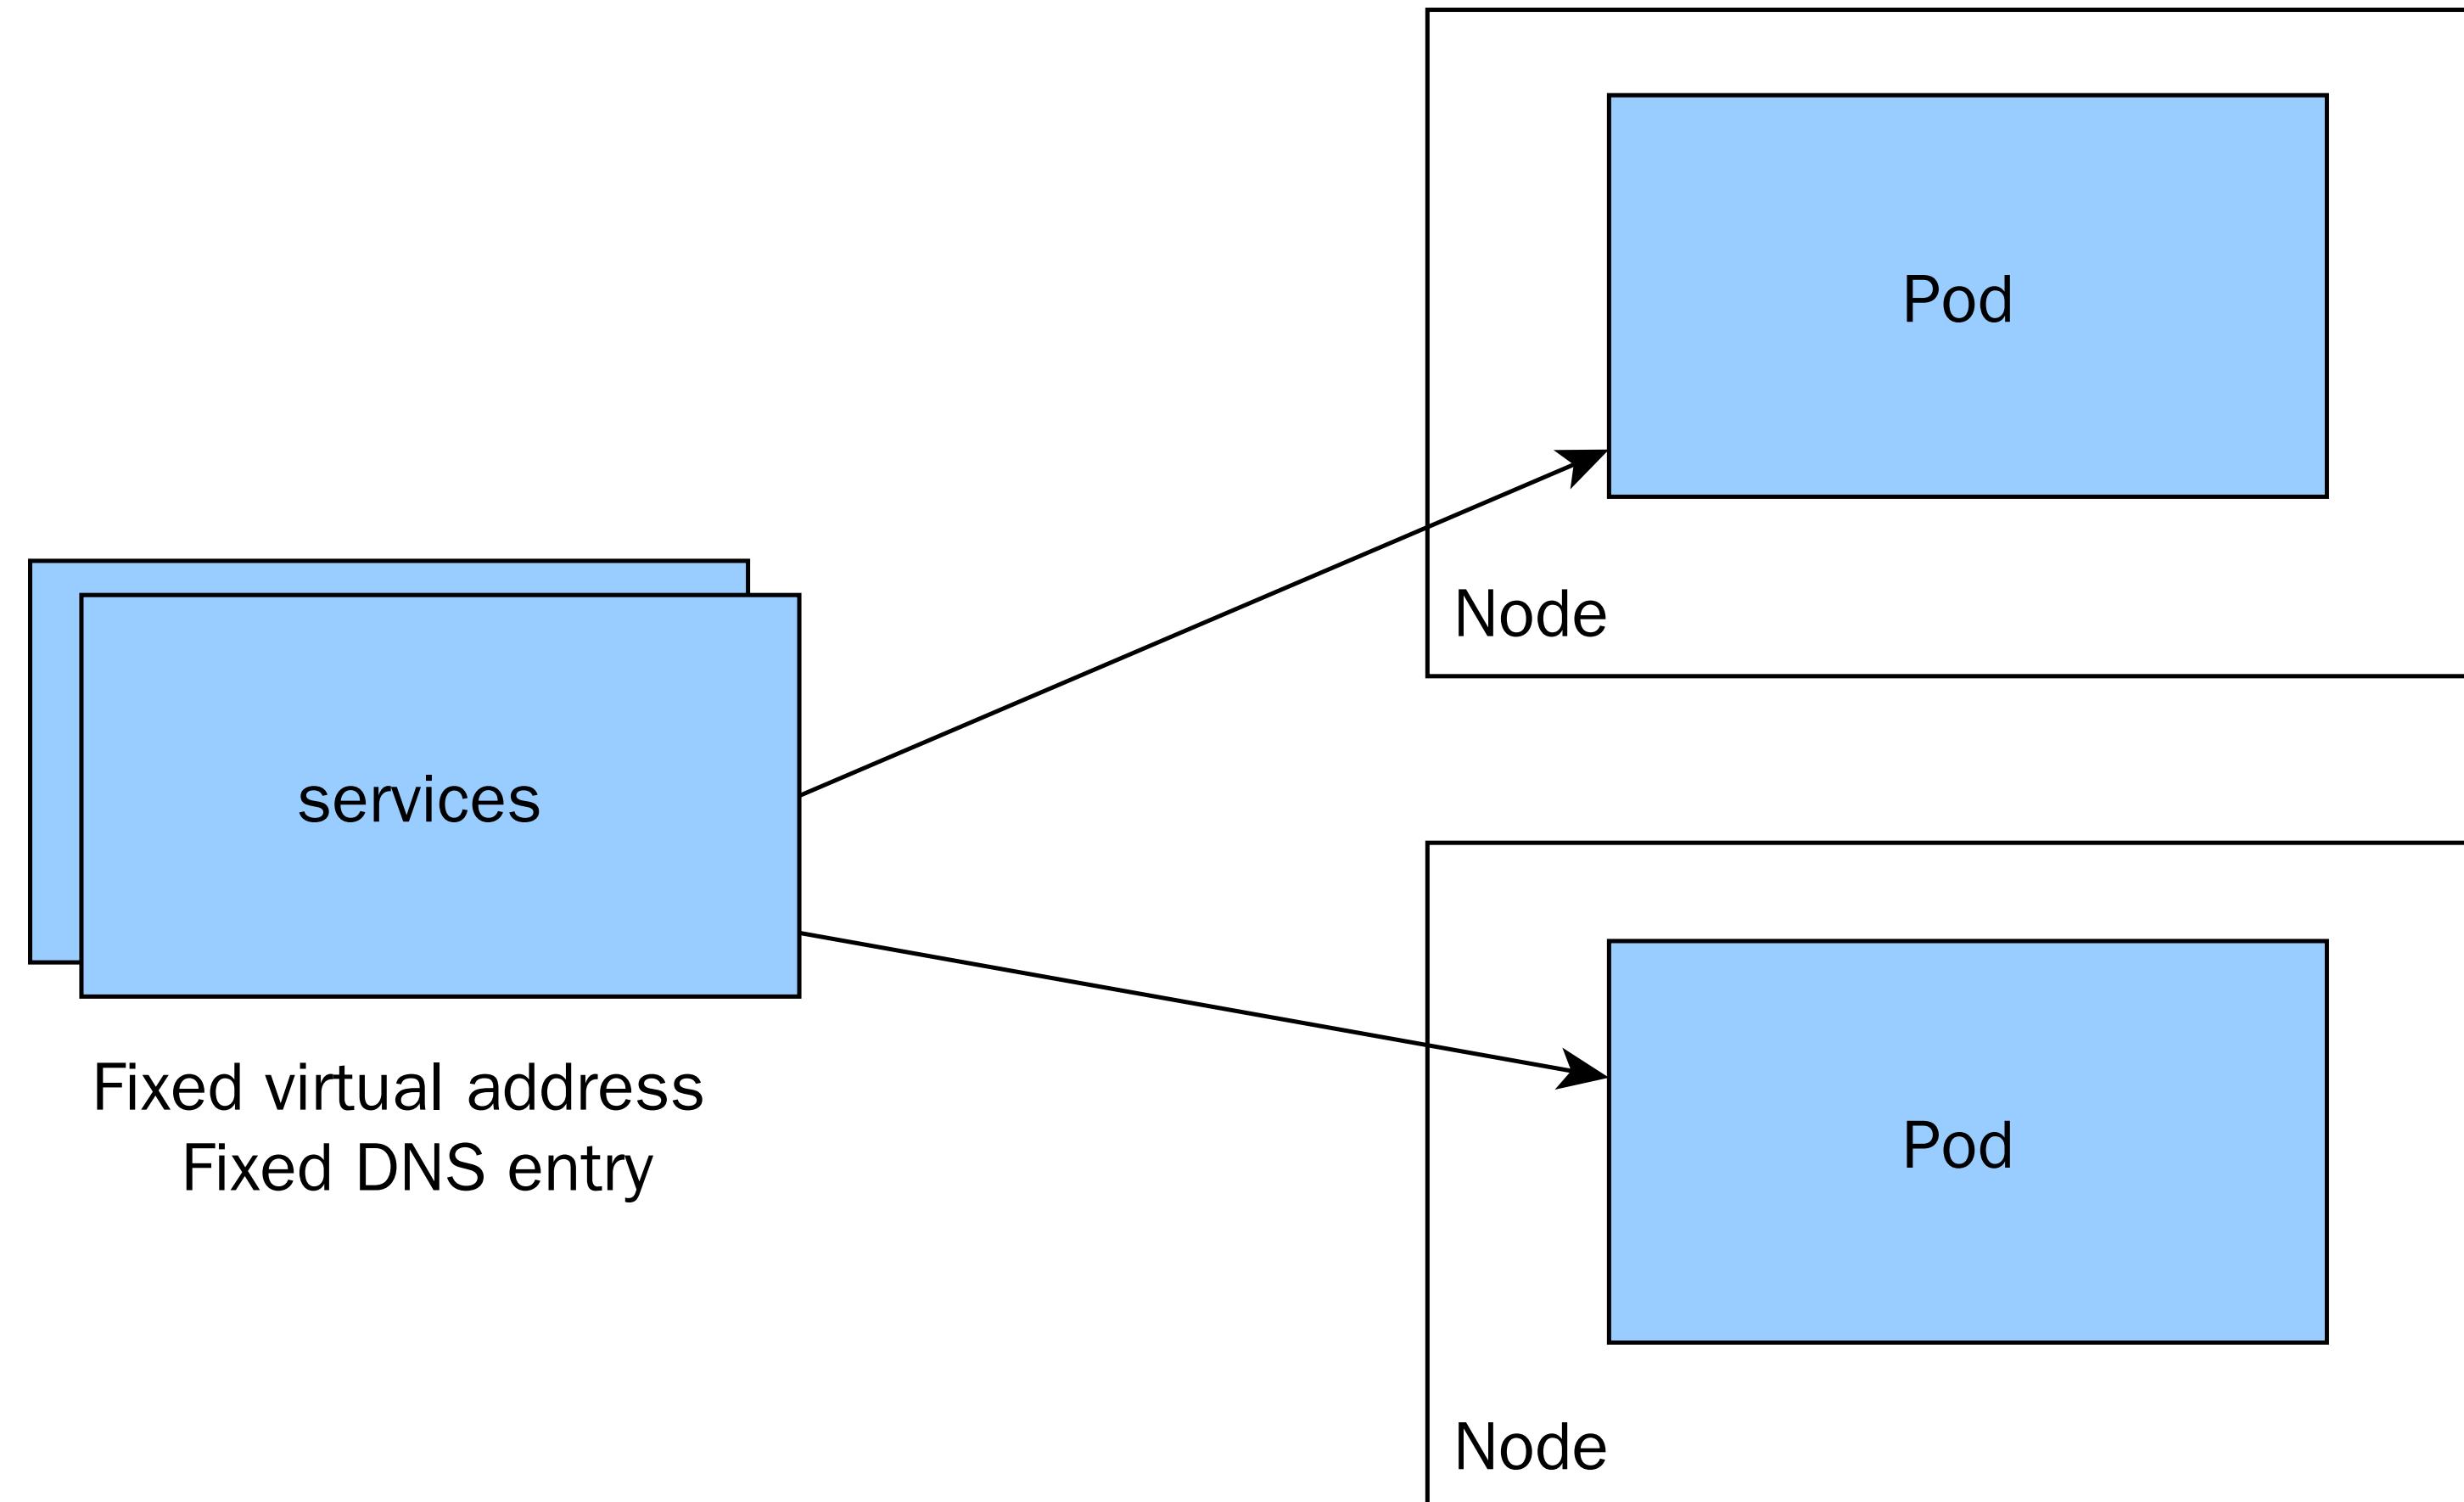
## Basic concepts

Name	Purpose	
Service	Interface	Entry point (Service Name)
Deployment	Factory	How many pods, which pods
Pod	Implementation	1+ docker running

# Deployment and pods



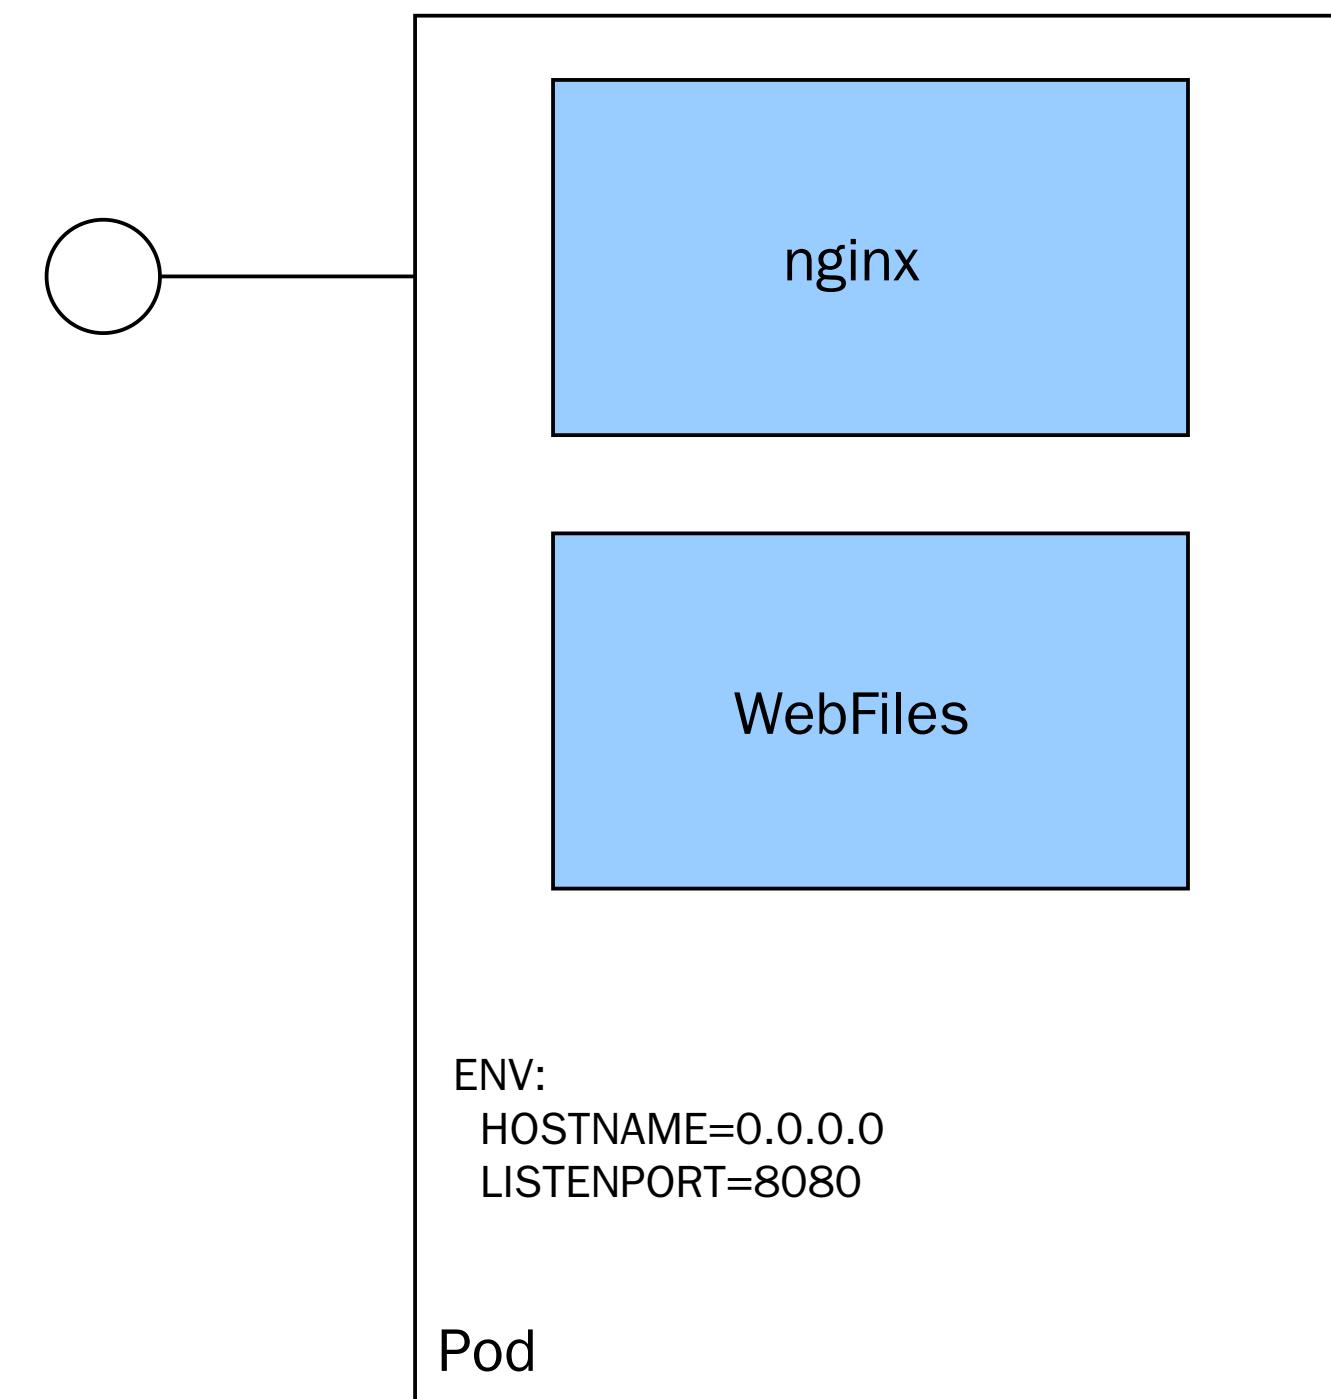
# Service and pods



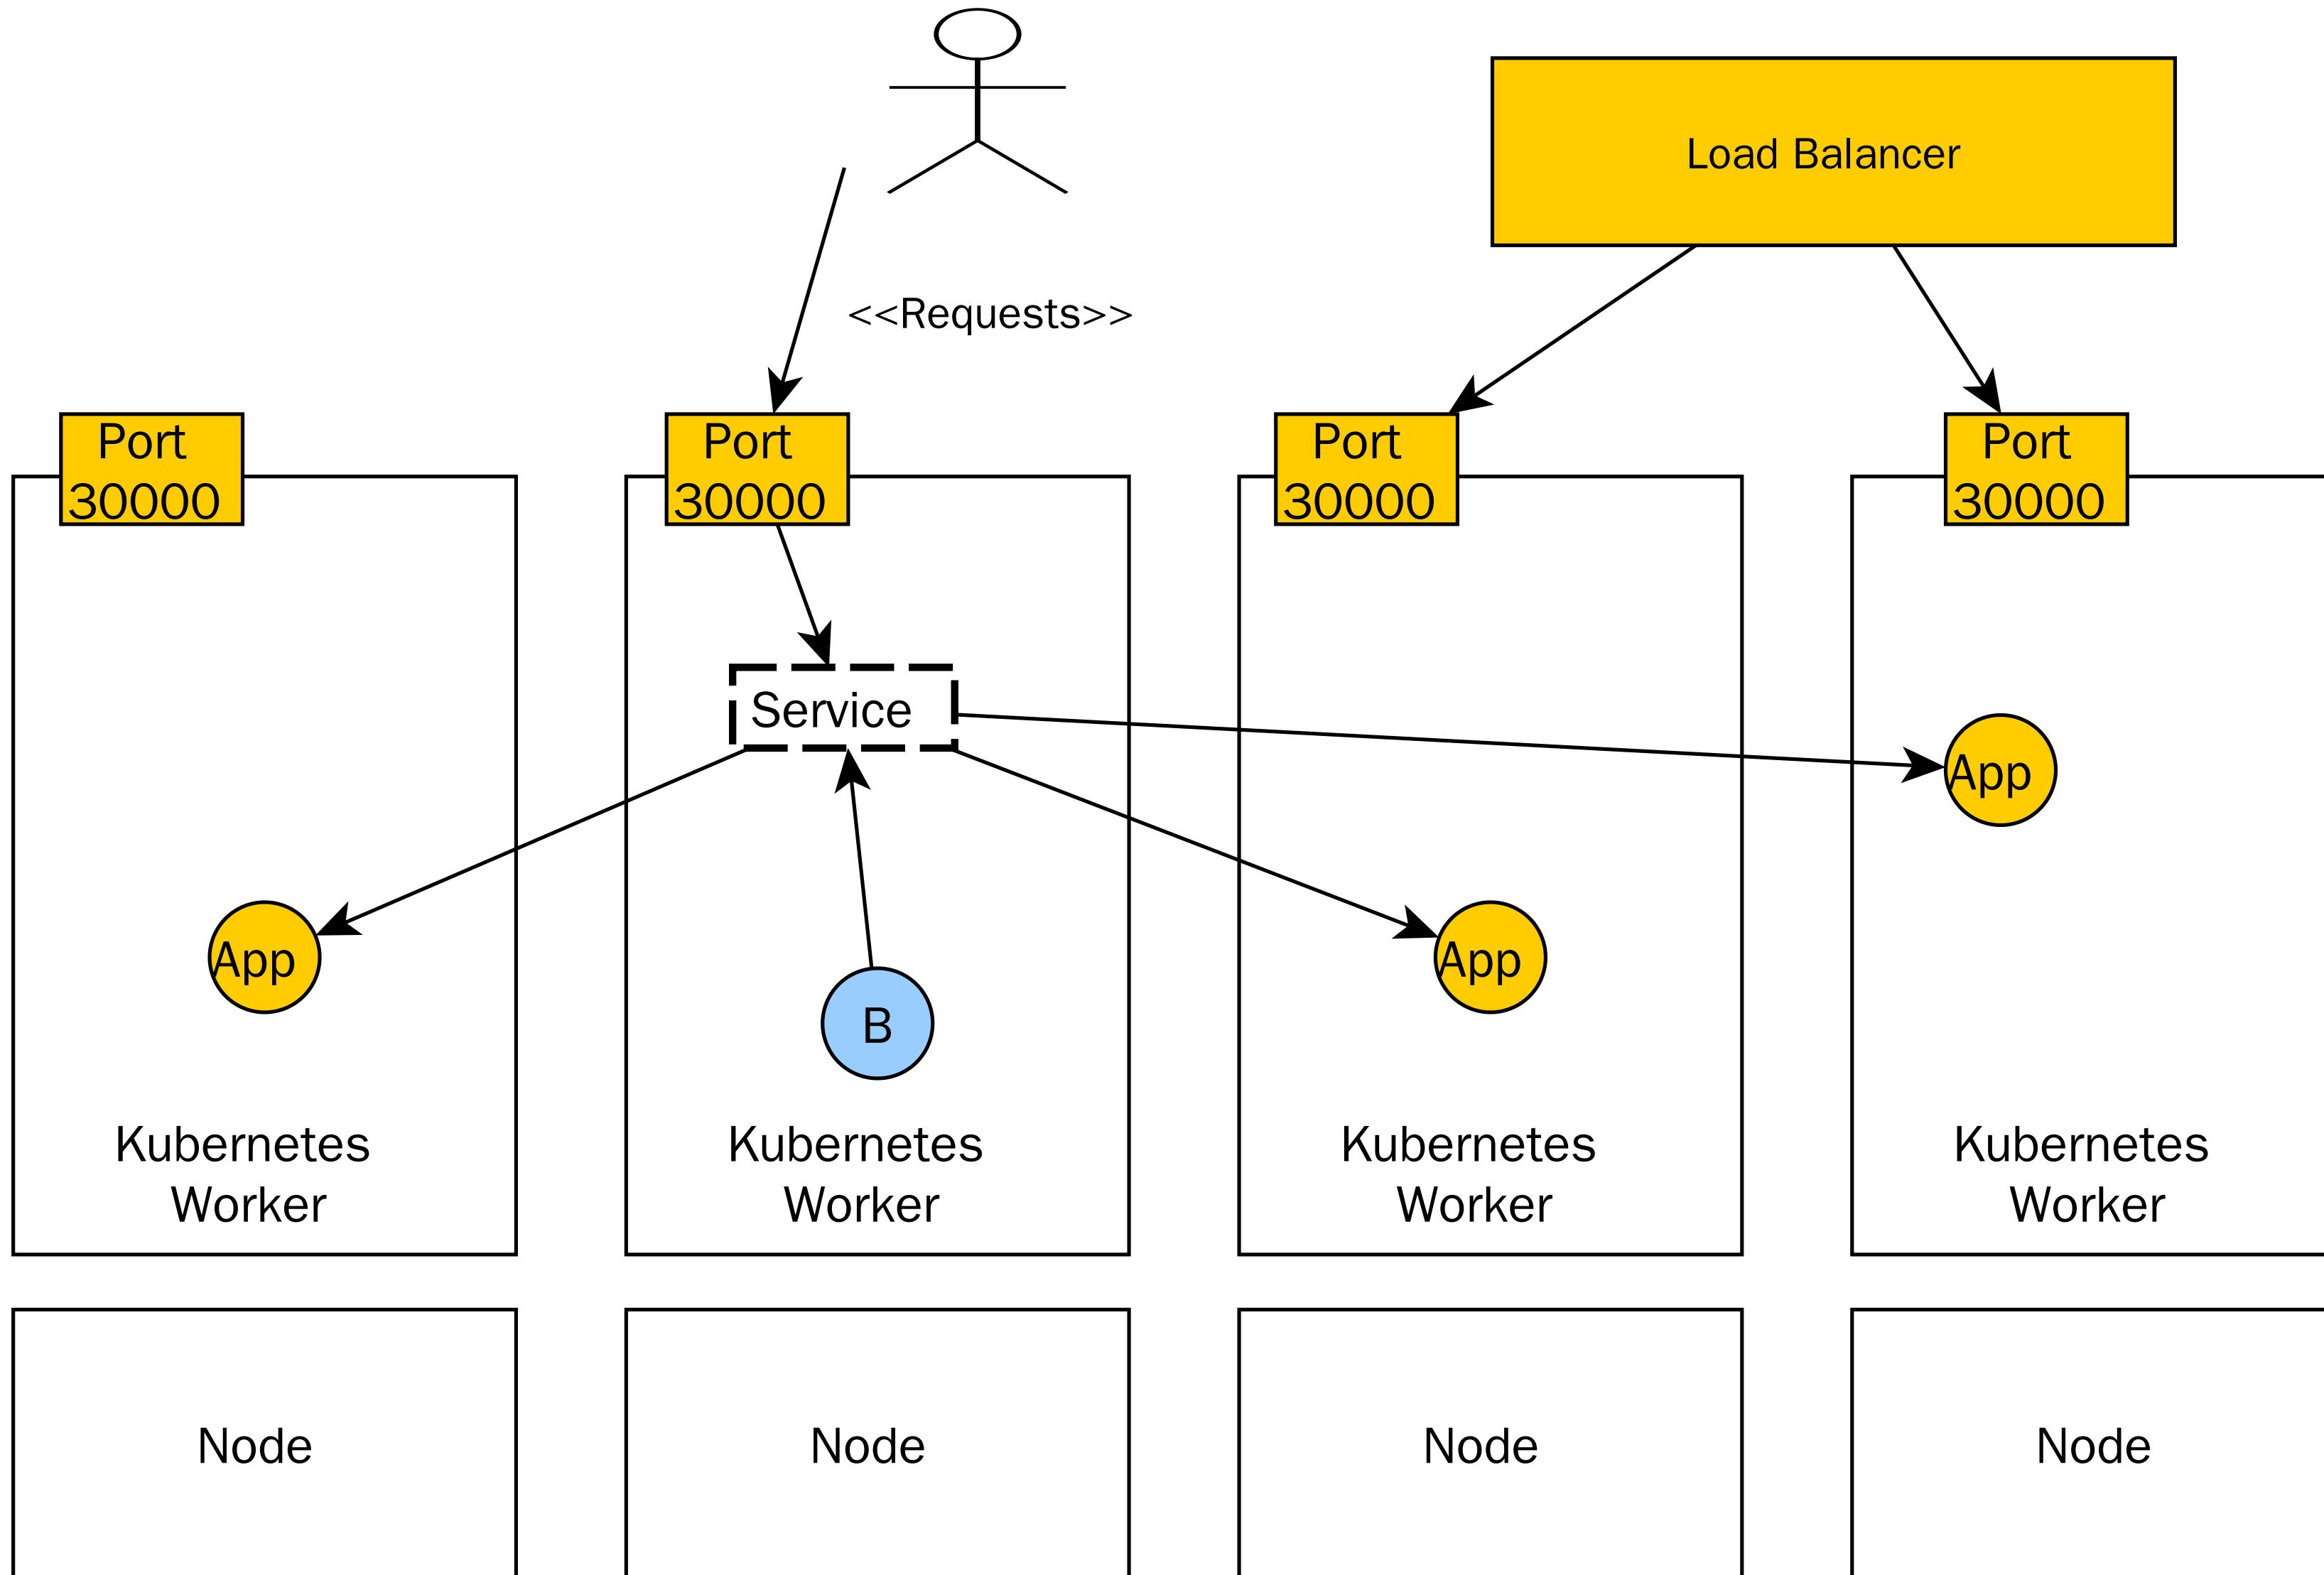
Service matches pods based on labels

# Pods

- See each other on localhost
- Live and die together
- Can expose multiple ports



# Load-balancing



# Service Discovery and labeling

- names in DNS:

curl http://users/list

- labels:

name=value

- annotations:

prometheus.io/scrape: "true"

## Service discovery

- loosely couple components
- auto-wiring with logging and monitoring
- drop-in installation (traefik, prometheus, ...)

## All hail the git!

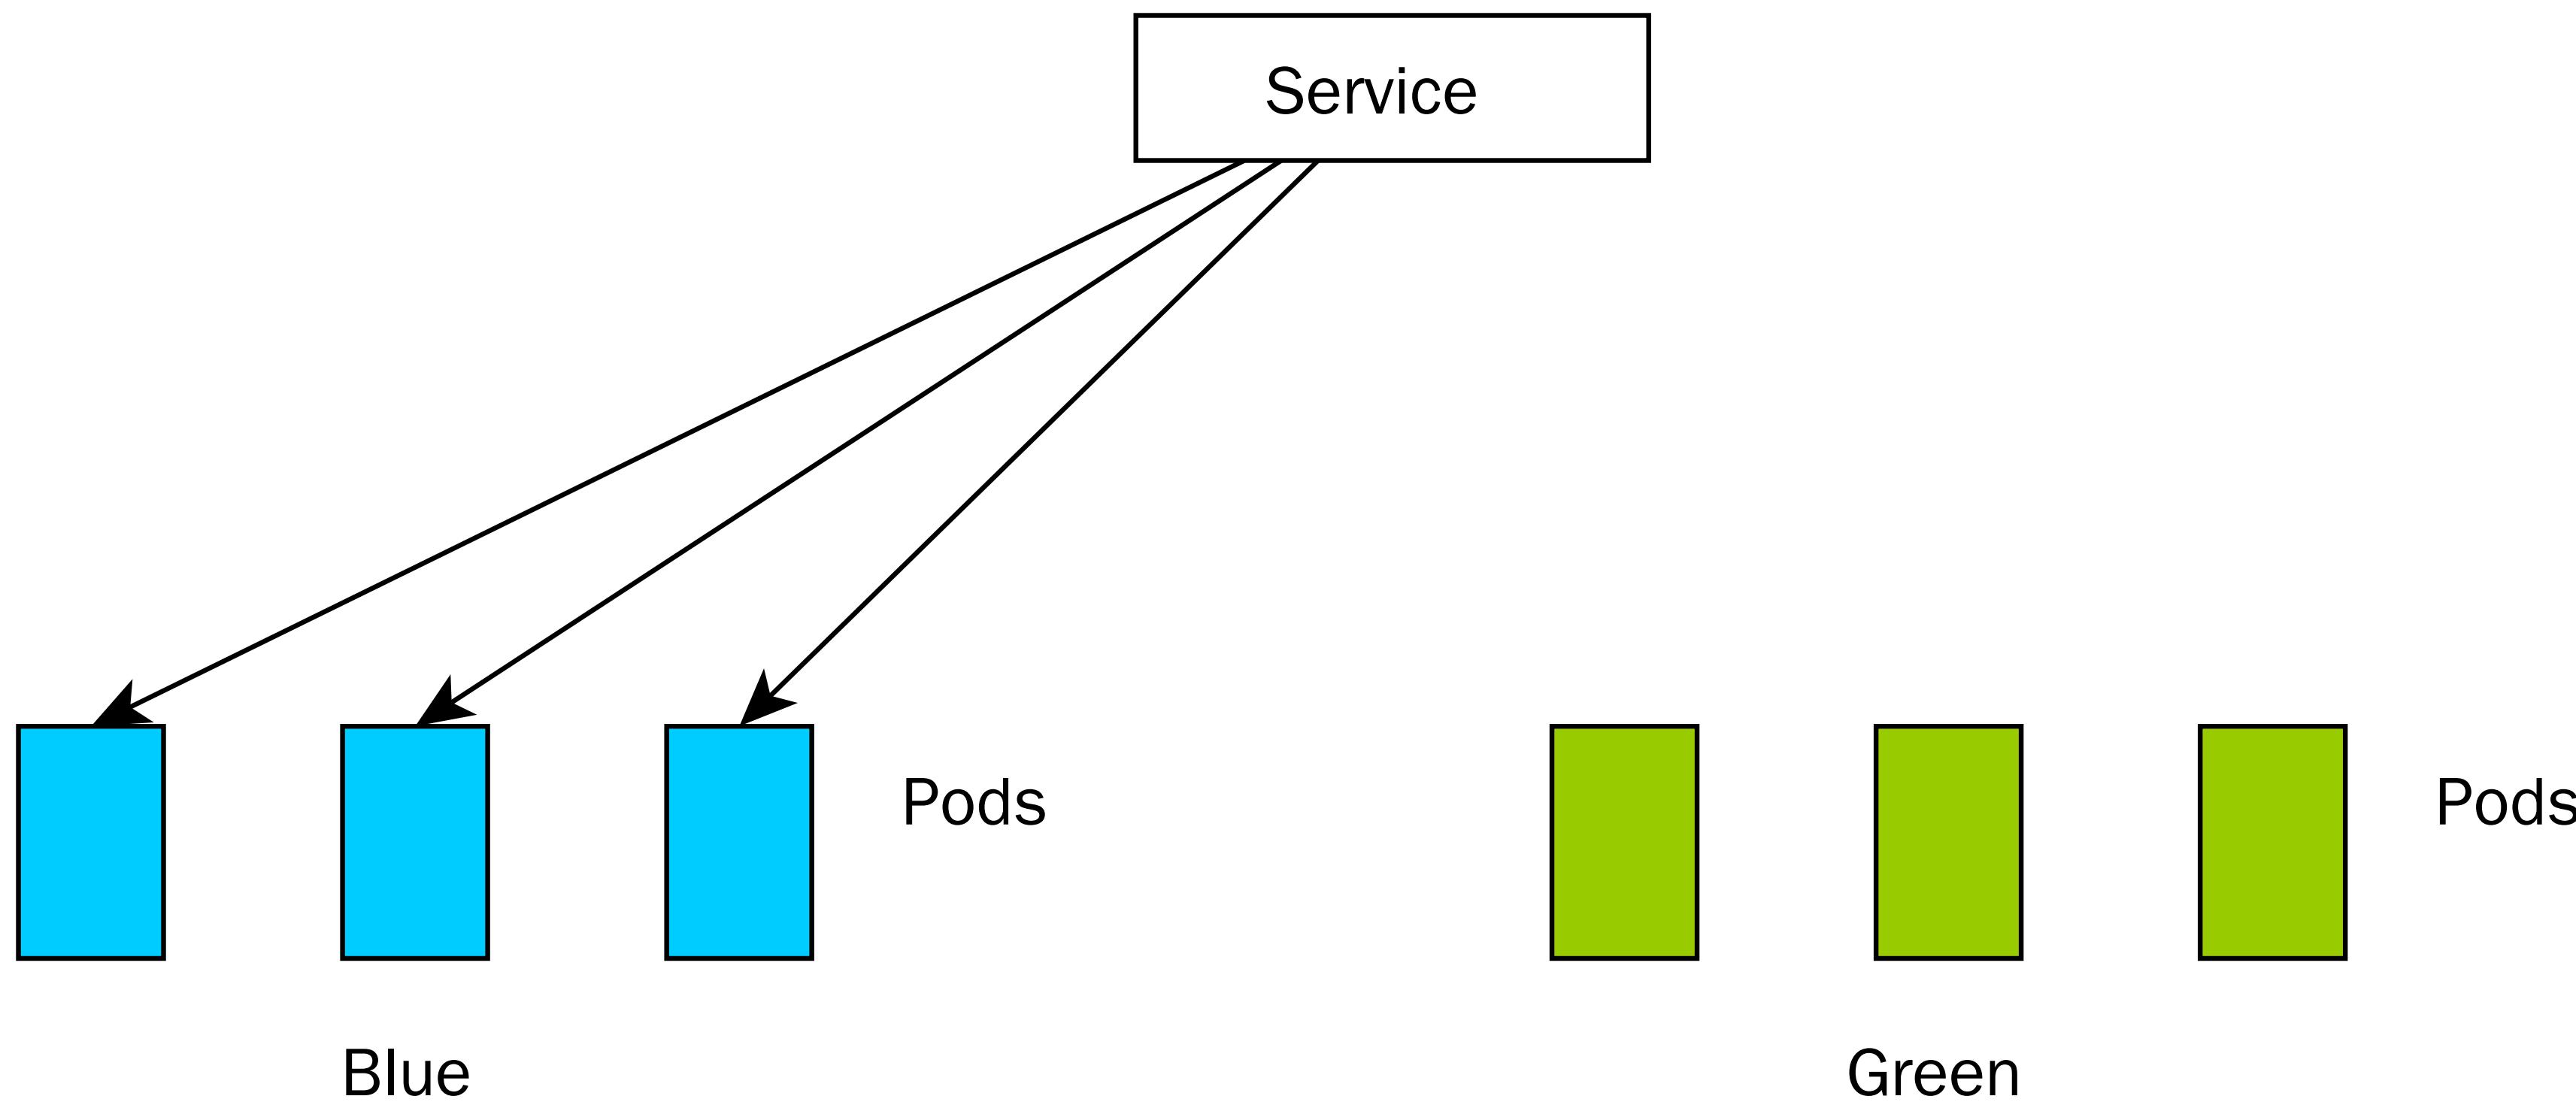
- Yaml
- integration:
  - monitoring, alarming, ingress-controller
- ...
- Devs can forget about infra... almost
- DevOps Culture Dream!

## Advanced!

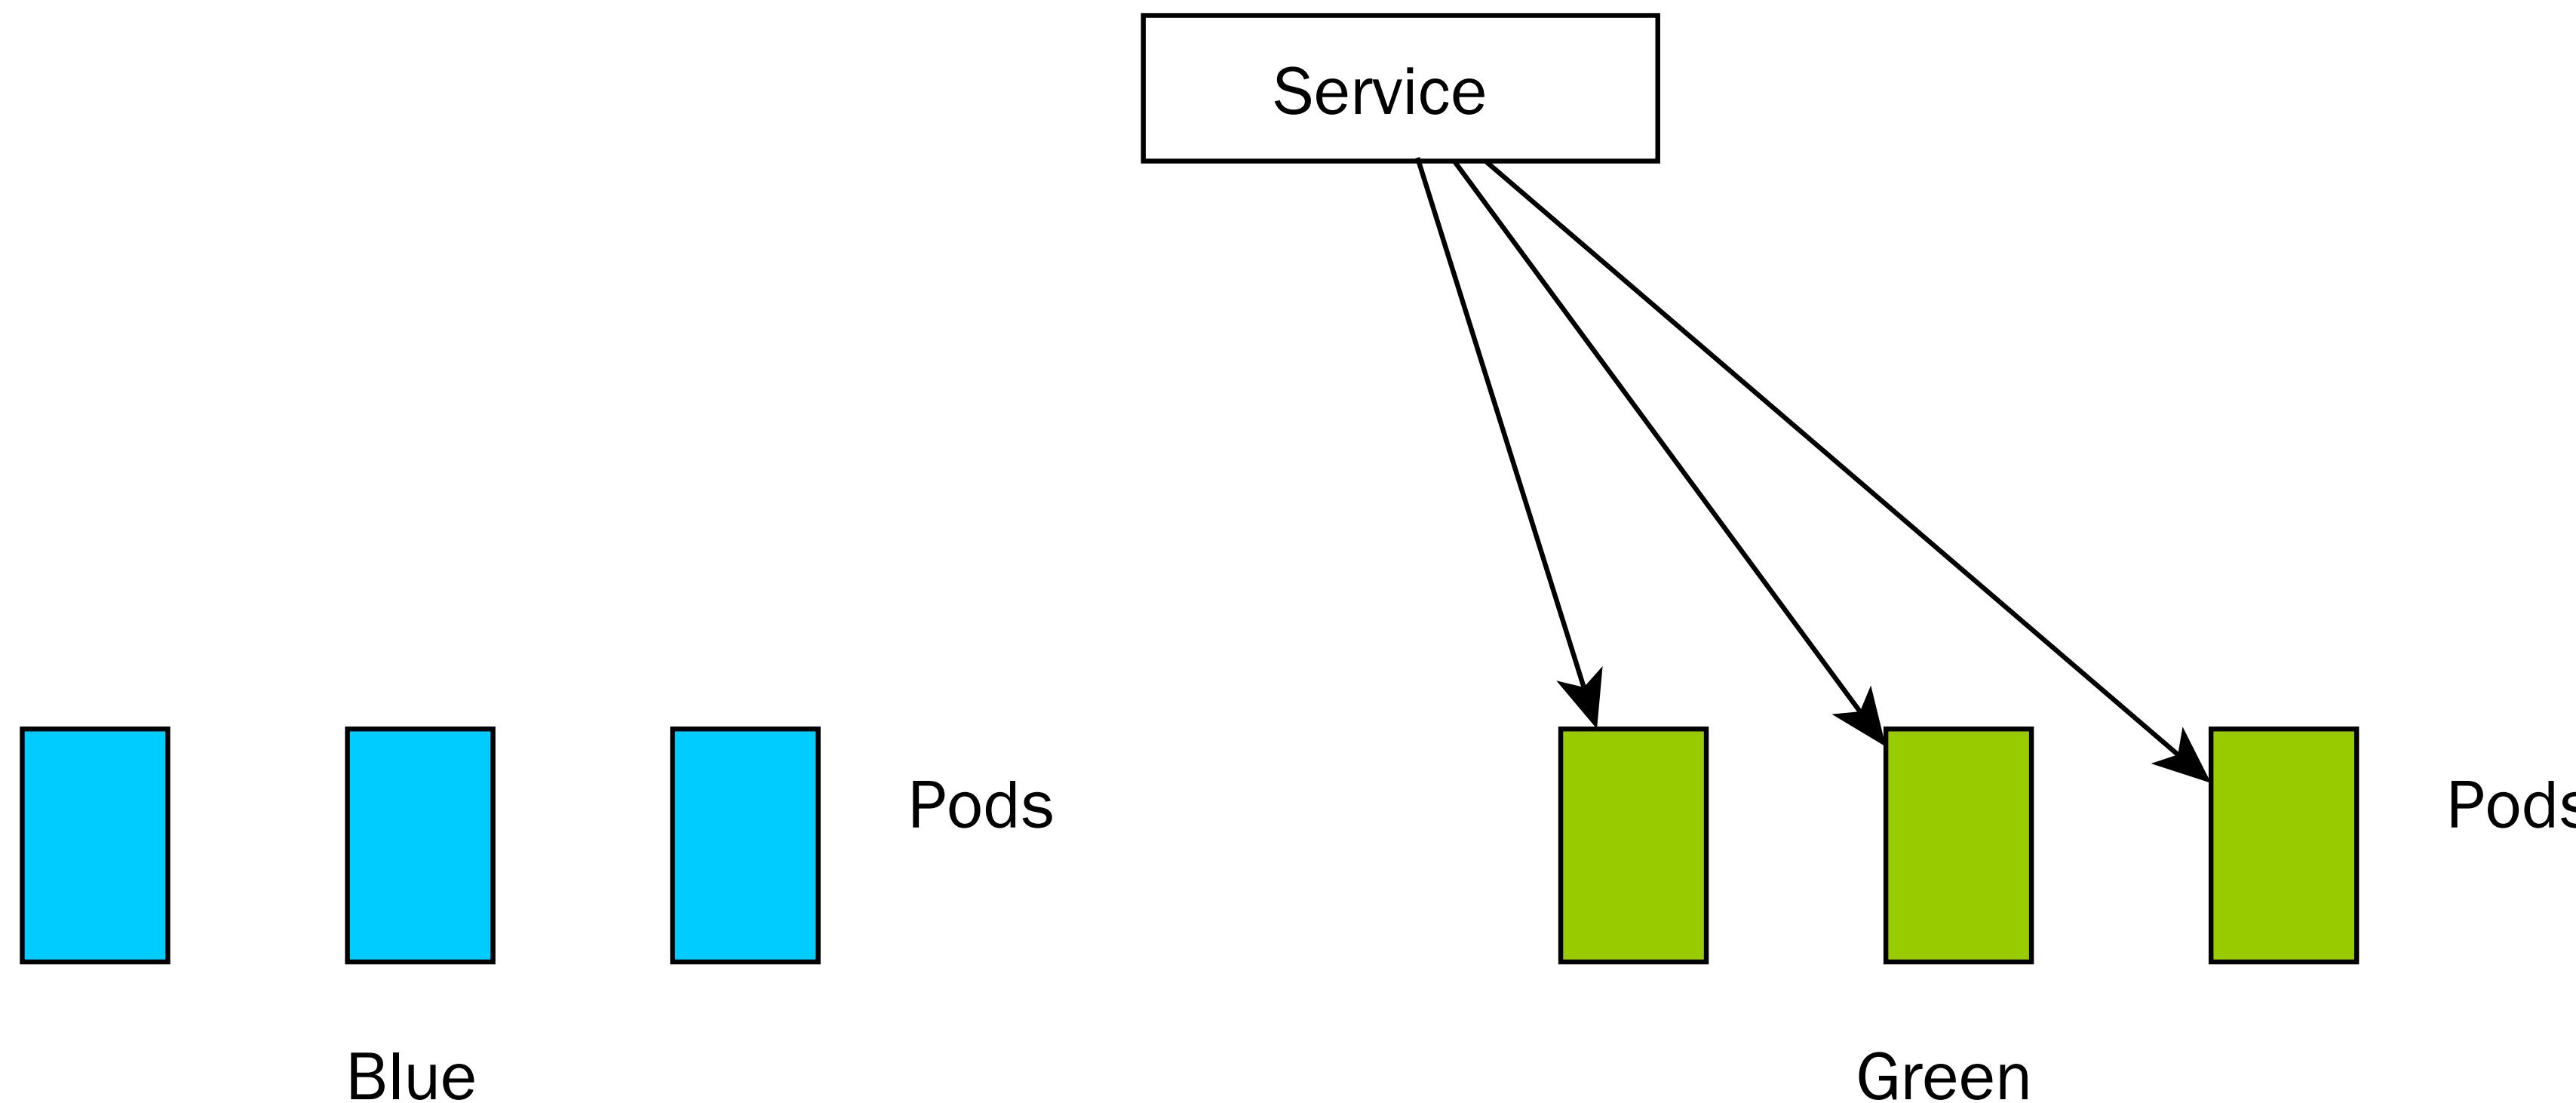
Sophisicated deployment strategies:

- rollout
- blue/green
- canary
- zero downtime deployments

# DEMO - GREEN/BLUE



# DEMO - GREEN/BLUE



## DEMO - GREEN/BLUE

```
kubectl patch service api-status \  
-p '{ "spec": { "selector": { "label": "green" } } }'
```

## DEMO - GREEN/BLUE

```
kubectl patch service api-status \  
-p '{ "spec": { "selector": { "label": "blue" } } }'
```

## **Effective Kubernetes**

- Start with small iterations
- Learn as-you-go
- Keep Kubernetes Understable

## Effective Kubernetes

- Move configuration to runtime
- Do not terrorize your devs with K8S
- No free lunch... app must be smarter

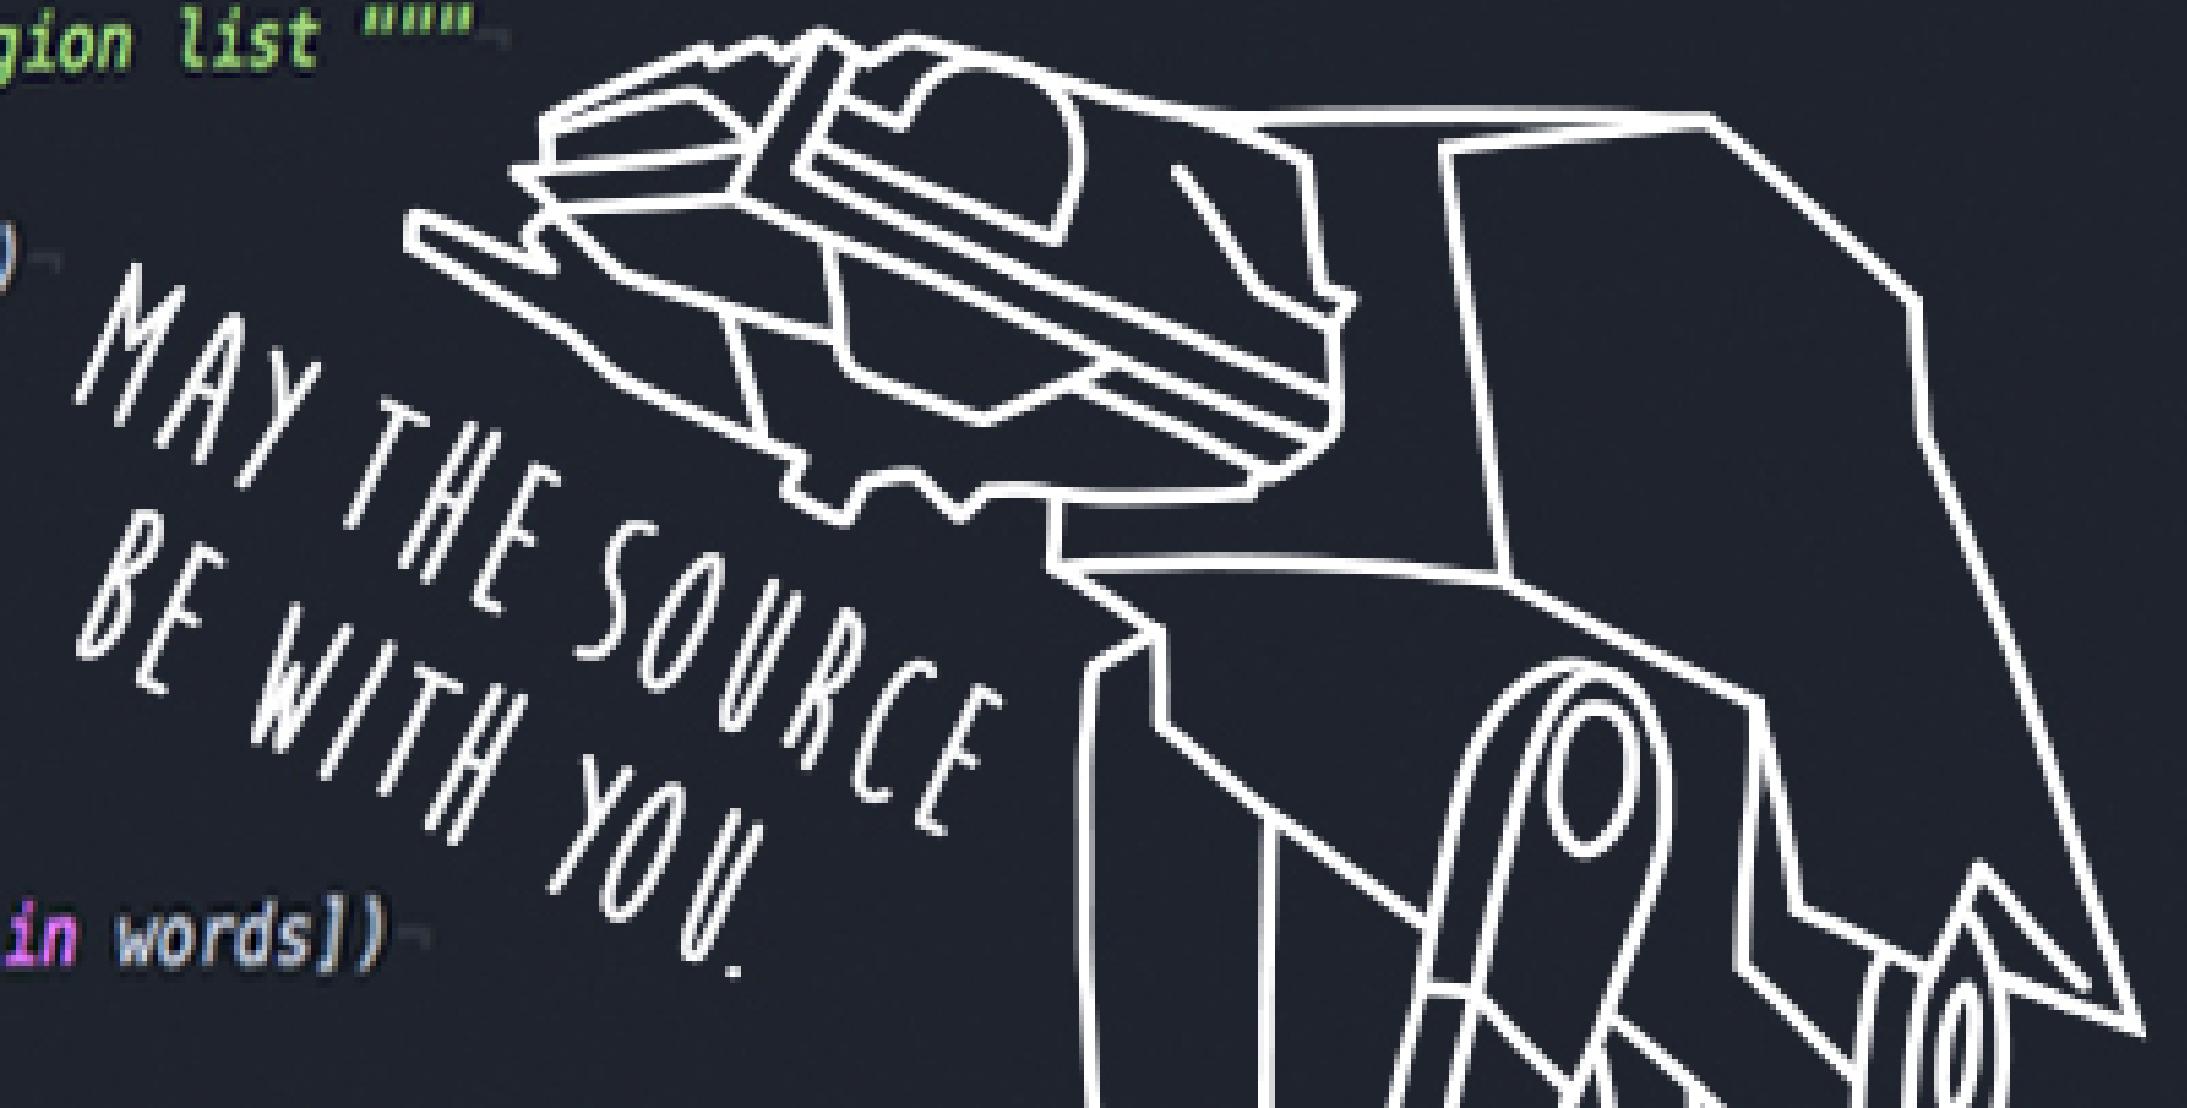
# THANK YOU. QUESTIONS?

```
123 def distance_matrix(regions):
124     """ Computes a distance matrix against a region list """
125     tuples = [r.as_tuple() for r in regions]
126     return cdist(tuples, tuples, region_distance)
127
128
129 def clusterize(words, **kwargs):
130     # TODO: write a cool docstring here
131     db = DBSCAN(metric="precomputed", **kwargs)
132     X = distance_matrix([Region.from_word(w) for w in words])
133     labels = [int(l) for l in db.fit_predict(X)]
```



[https://github.com/wojciech12/workshop\\_kubernetes\\_and\\_cloudnative](https://github.com/wojciech12/workshop_kubernetes_and_cloudnative)

```
123 def distance_matrix(regions):
124     """ Computes a distance matrix against a region list """
125     tuples = [r.as_tuple() for r in regions]
126     return cdist(tuples, tuples, region_distance)
127
128
129 def clusterize(words, **kwargs):
130     # TODO: write a cool docstring here
131     db = DBSCAN(metric="precomputed", **kwargs)
132     X = distance_matrix([Region.from_word(w) for w in words])
133     labels = [int(l) for l in db.fit_predict(X)]
```



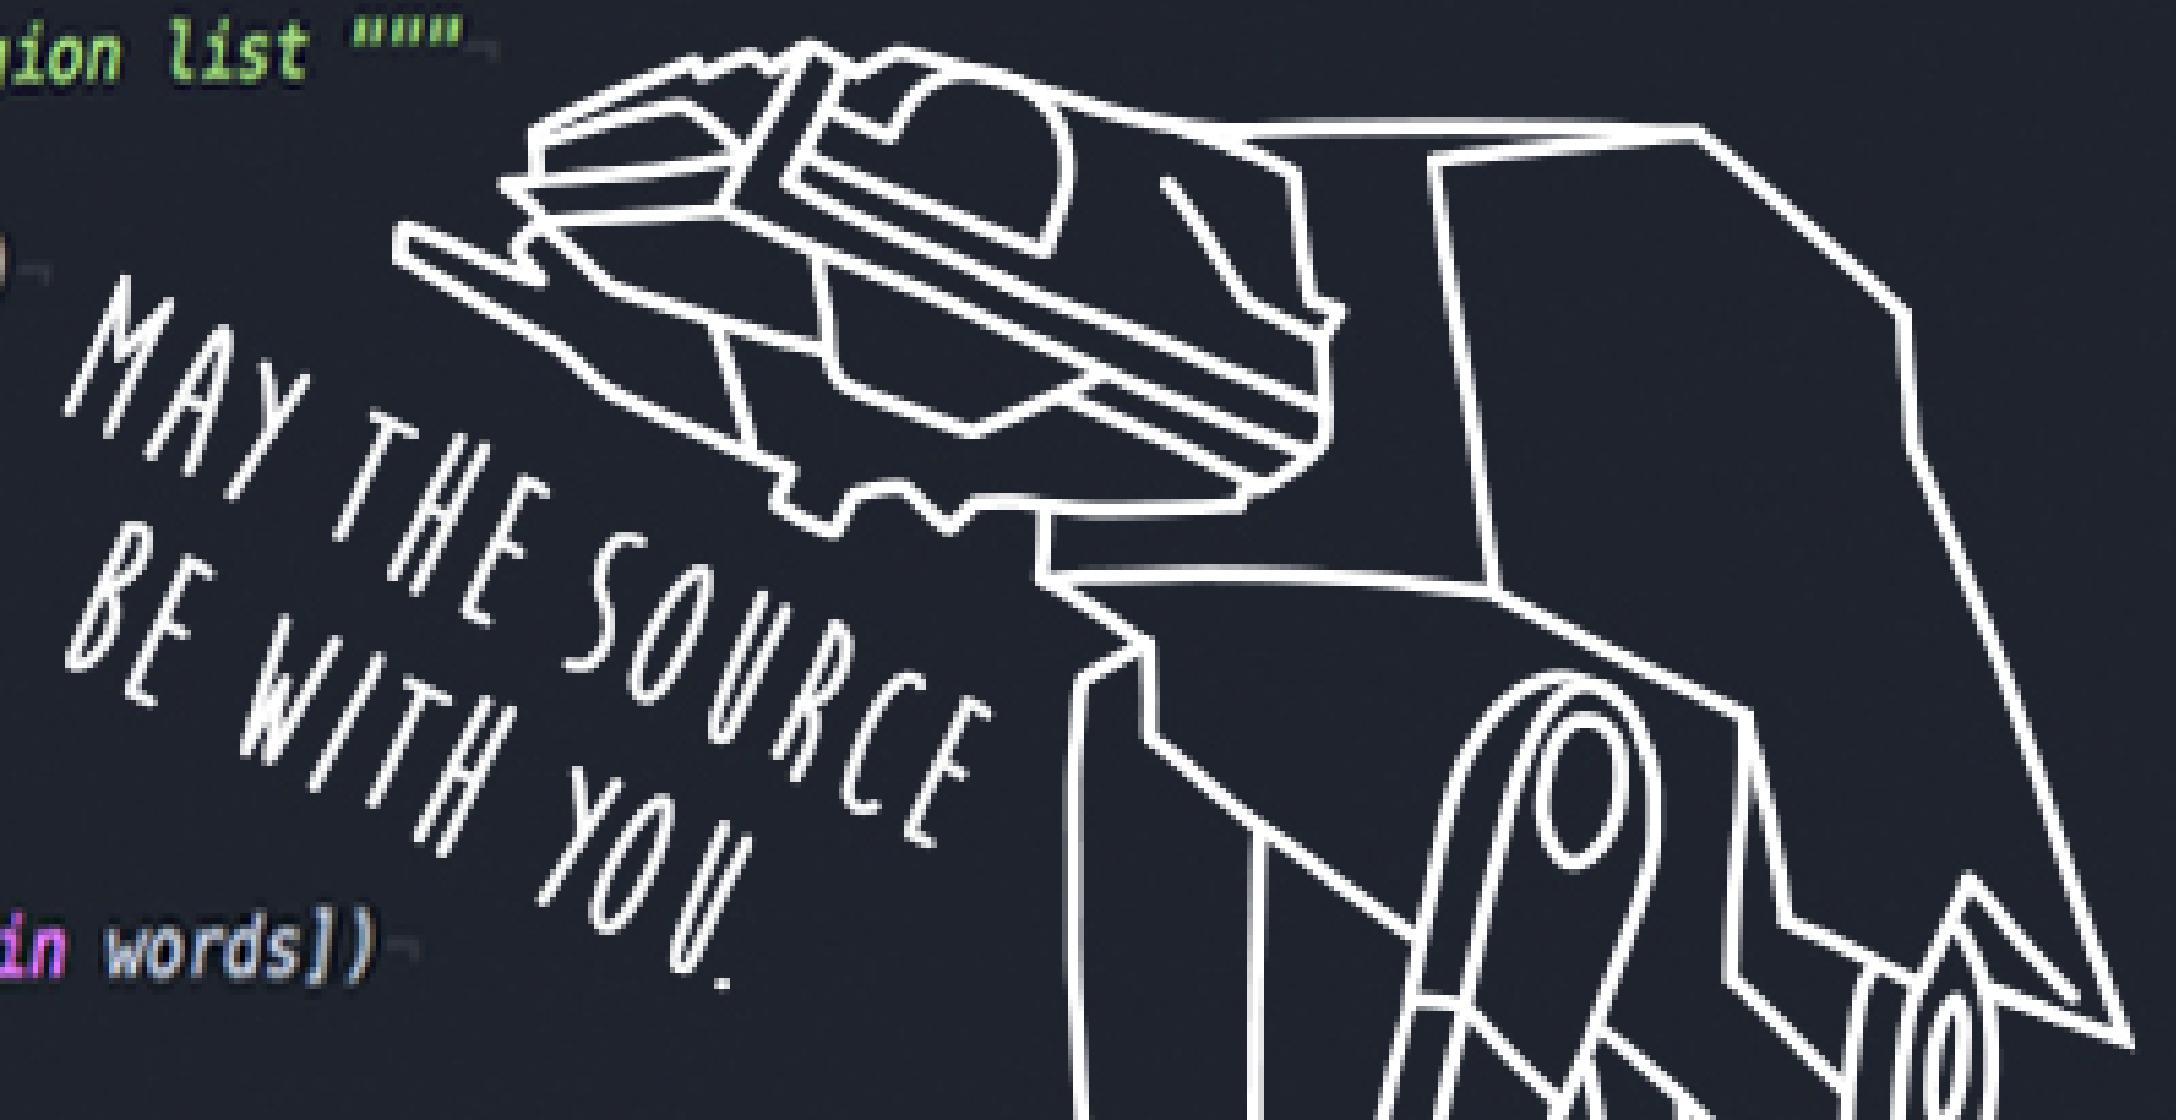
[https://github.com/wojciech12/talk\\_zero\\_downtime\\_deployment\\_with\\_kubernetes](https://github.com/wojciech12/talk_zero_downtime_deployment_with_kubernetes)

```
123 def distance_matrix(regions):
124     """ Computes a distance matrix against a region list """
125     tuples = [r.as_tuple() for r in regions]
126     return cdist(tuples, tuples, region_distance)
127
128
129 def clusterize(words, **kwargs):
130     # TODO: write a cool docstring here
131     db = DBSCAN(metric="precomputed", **kwargs)
132     X = distance_matrix([Region.from_word(w) for w in words])
133     labels = [int(l) for l in db.fit_predict(X)]
```



# BACKUP SLIDES

```
123 def distance_matrix(regions):
124     """ Computes a distance matrix against a region list """
125     tuples = [r.as_tuple() for r in regions]
126     return cdist(tuples, tuples, region_distance)
127
128
129 def clusterize(words, **kwargs):
130     # TODO: write a cool docstring here
131     db = DBSCAN(metric="precomputed", **kwargs)
132     X = distance_matrix([Region.from_word(w) for w in words])
133     labels = [int(l) for l in db.fit_predict(X)]
```



# INFRA TOOLS

- Prometheus + AlertMnager + Grafana
- Traefik
- Kafka - Yolean/kubernetes-kafka
- Vault on Etcd - banzaicloud/bank-vaults

# DATA STORES

- Kafka - Yolean/kubernetes-kafka
- Etcd - coreos/etcd-operator
- DB: PSQL and Mongo

## BACKUPS:

- old-school backups with ARK and Restic
- replications across clouds

# 1. CLEAN UP

- Single script for repo - Makefile [1]
- Resurrect the README

[1] With zsh or bash auto-completion plug-in in your terminal.

## 2. GET BACK ALL THE KNOWLEDGE

Extract from:

- Puppet, ... ➔ Dockerfile
- Running Instances ➔ Dockerfile, README.rst
- Nagios, ... ➔ README.rst + `checks/`

## 3. INTRODUCE RUN\_LOCAL

- make run\_local
- A nice section on how to run in README.rst
- with docker-compose

The most crucial point.

## 4. GET TO KUBERNETES

- make kube\_create\_config
- make kube\_apply
- Generate the yaml files if your envs differ  
secrets from gopass (password manager)

## 5. CONTINUOUS DEPLOYMENT

Travis:

- the same Makefile as dev use

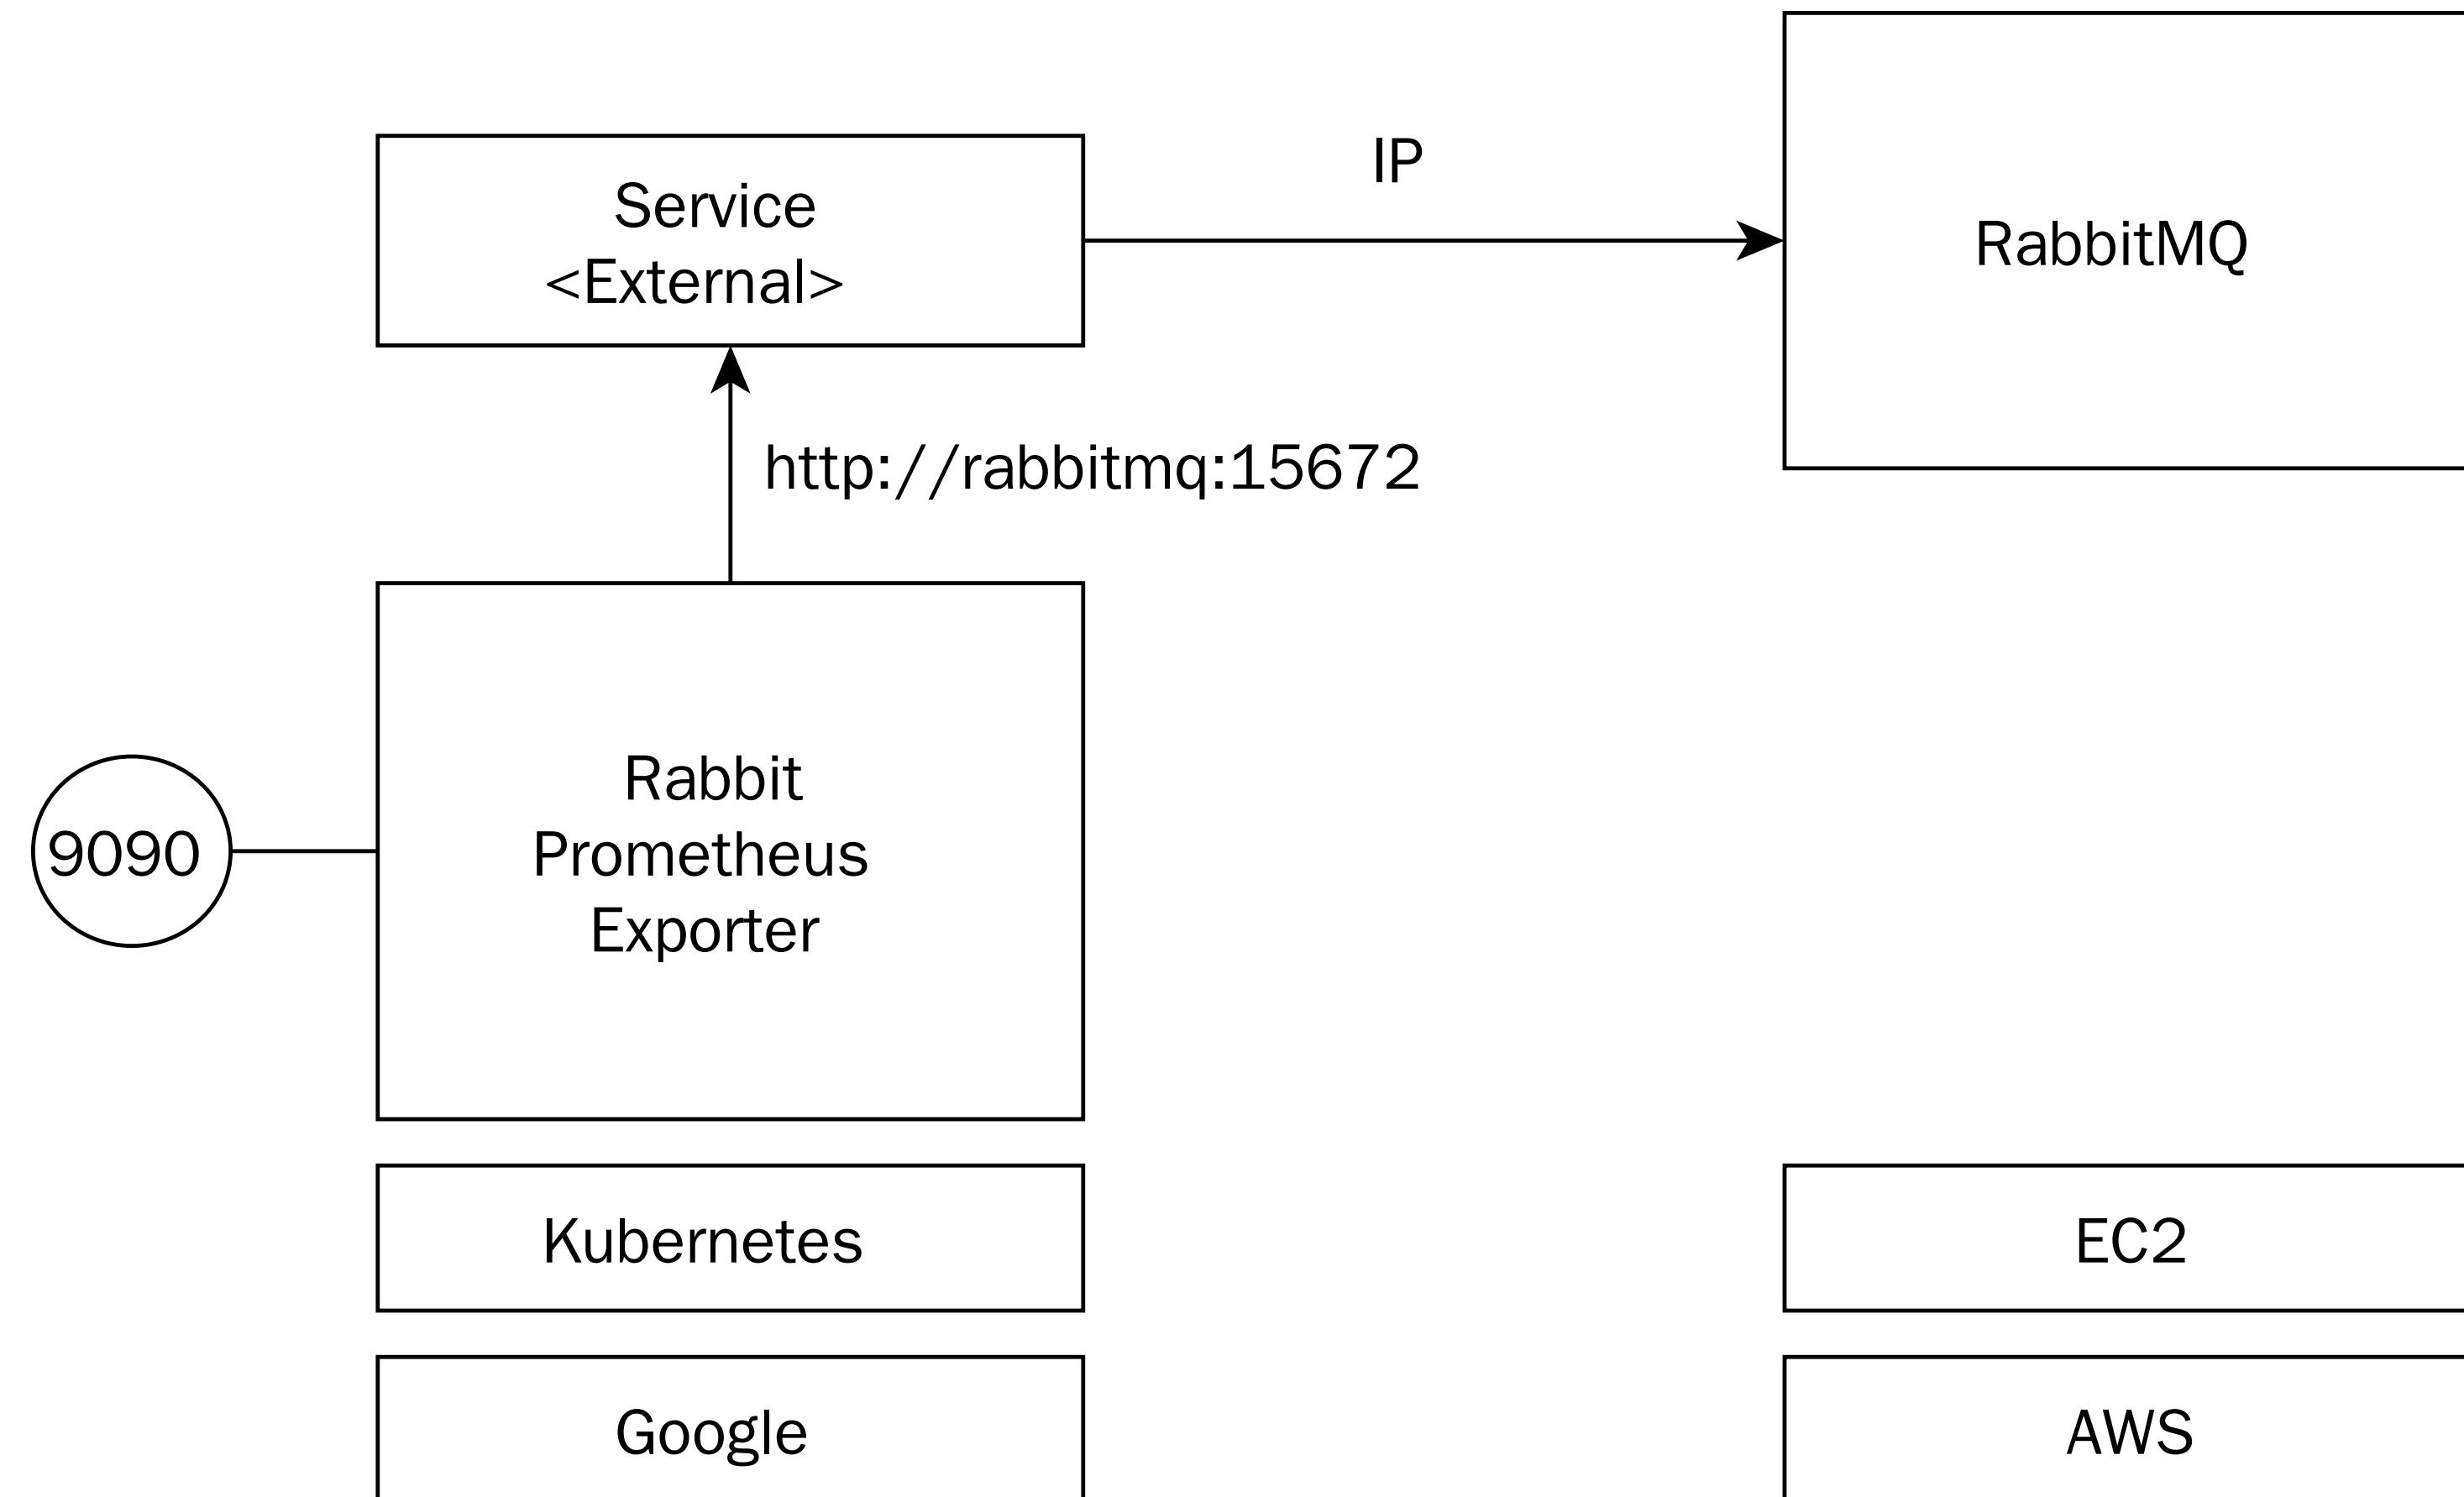
## 6. KEEP IT RUNNING

Bridge the new with old:

- Use External Services in Kubernetes
- Expose k8s in the Legacy [1]

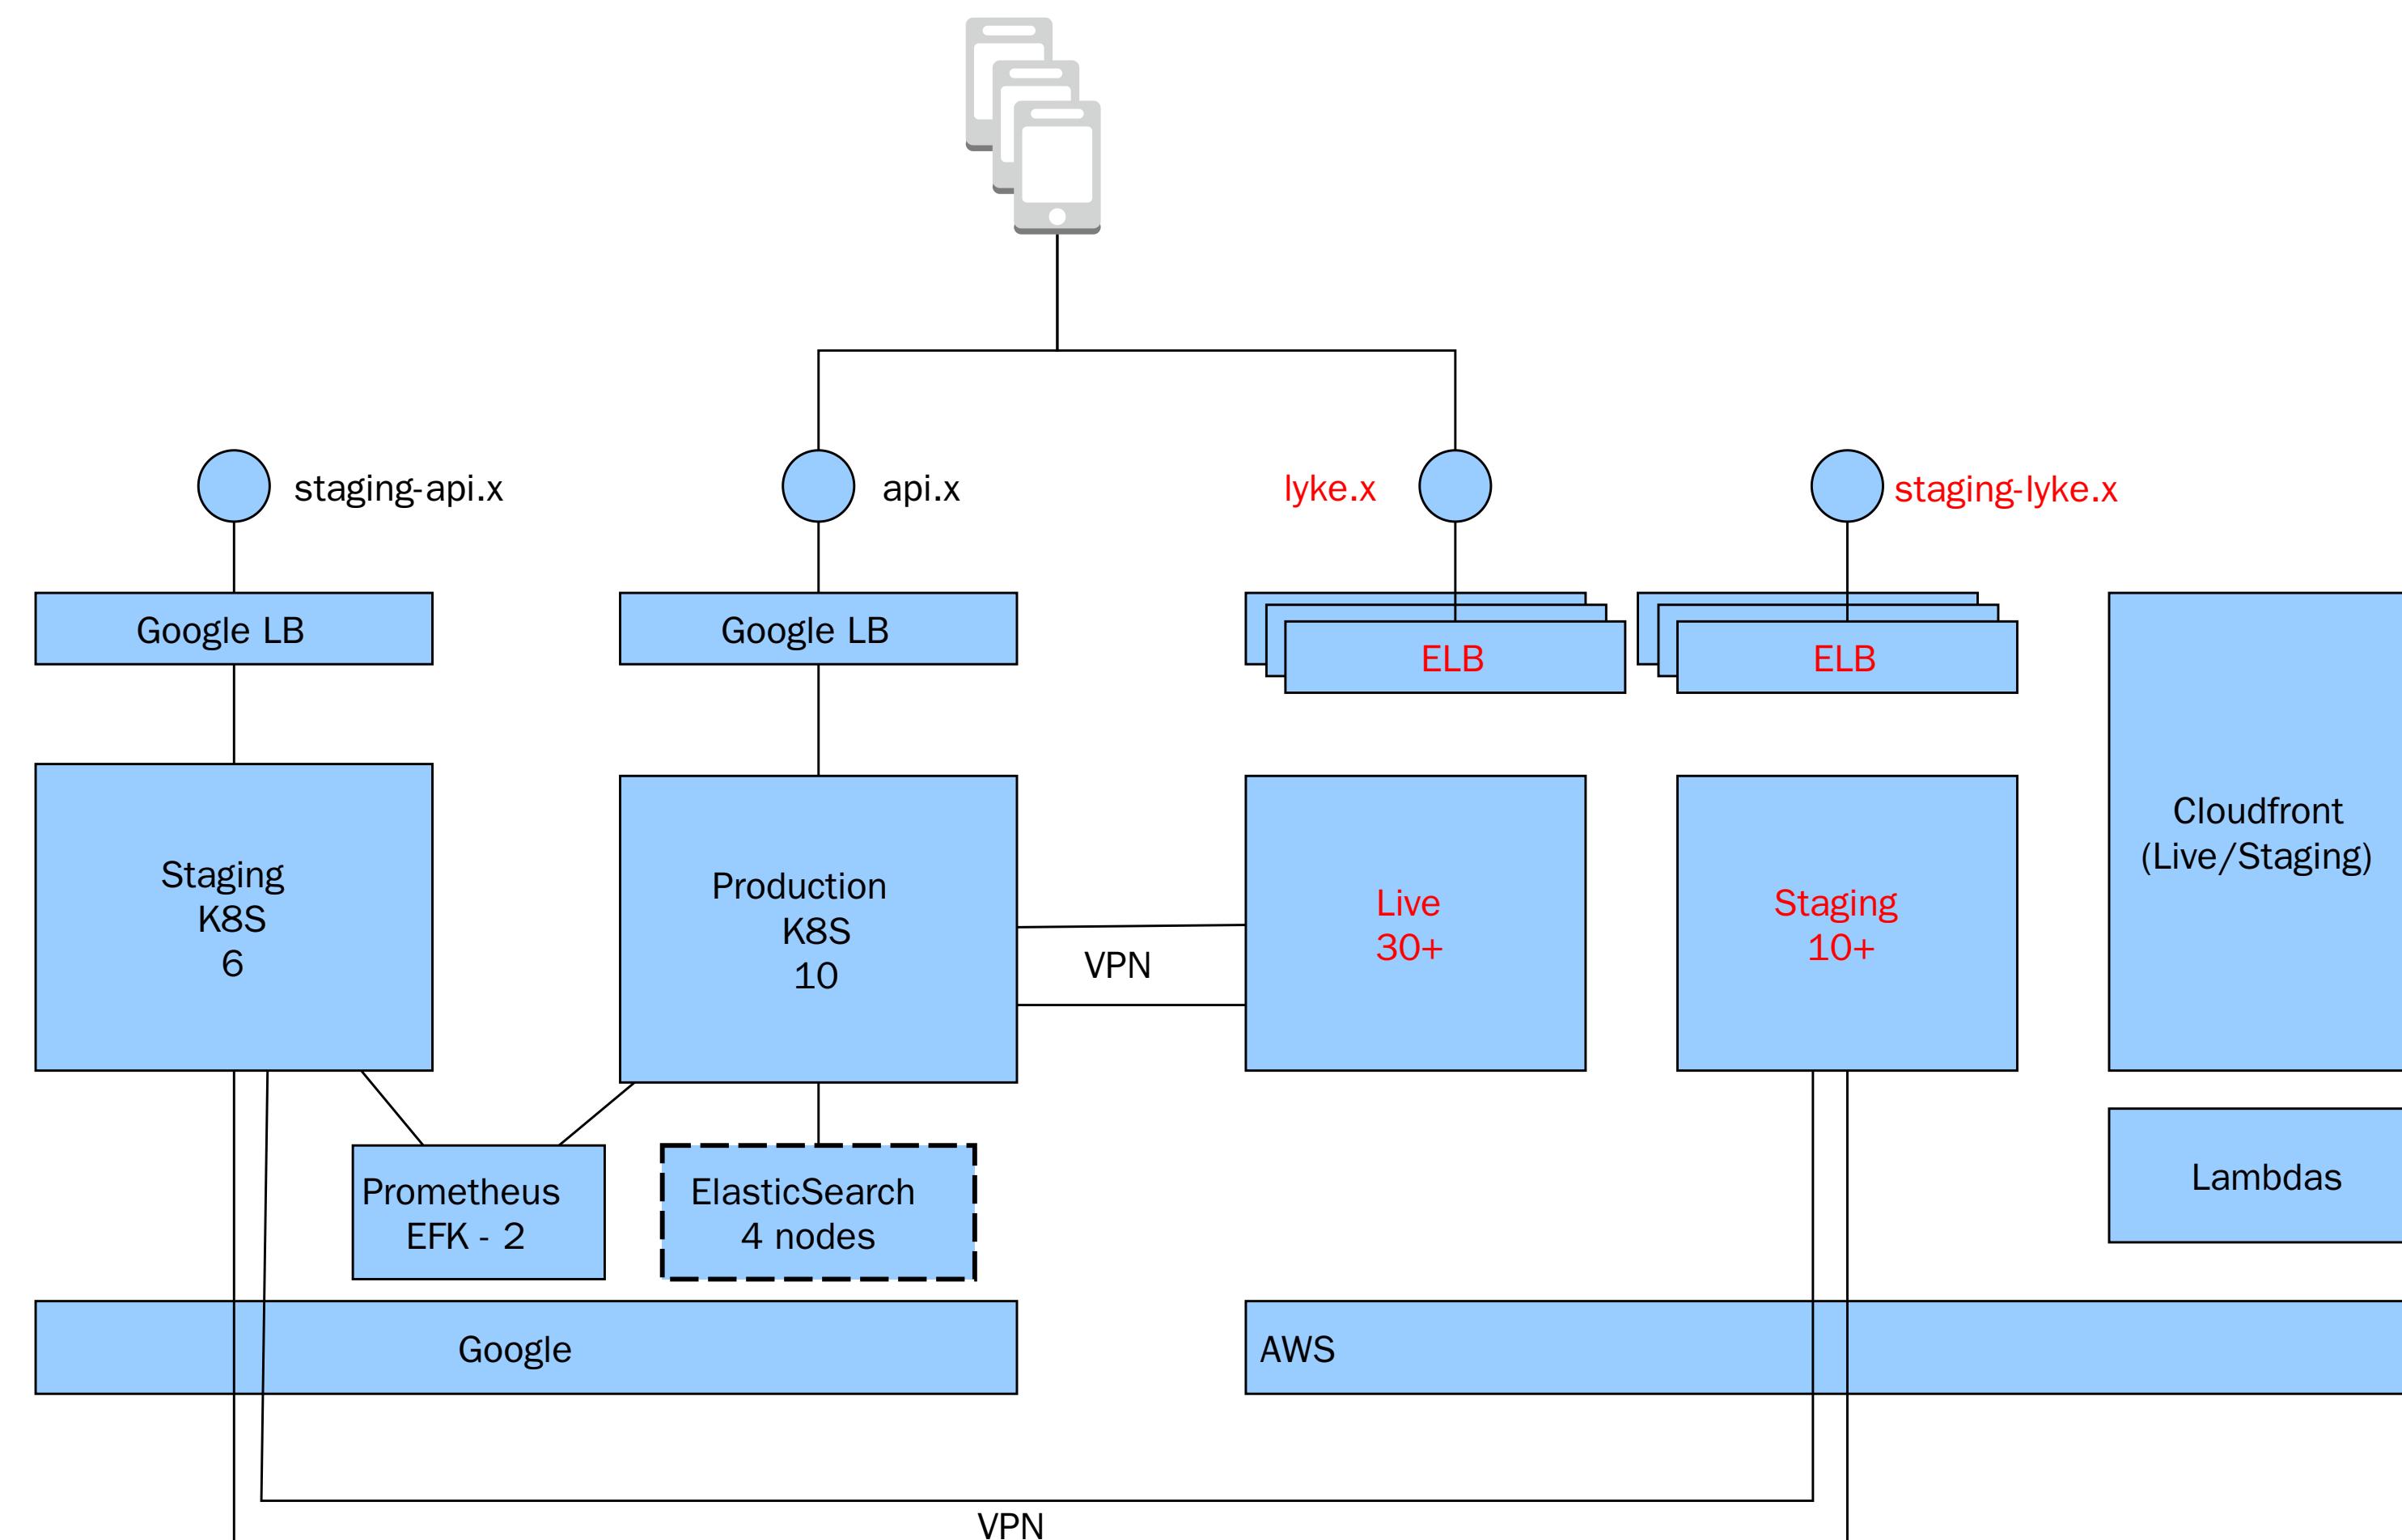
[1] hard coded IP:PORT, considered: K8S events to Consul

# Bridge the new with old



Monitor legacy with new stack

# Architecture During Migration



## 7. INTRODUCE SMOKE-TEST

```
TARGET_URL=127.0.0 make smoke_test  
TARGET_URL=api.example.com/users make smoke_test
```

## 8. SERVICE SELF-AWARE

Add to old services:

1. *metrics/*
2. *health/*
3. *info/*

## **9. MOVE TO MICRO-SERVICES**

Offload Legacy:

- Keep the lights on
- New functionality to micro-services

## 10. GET PERFORMANCE TESTING

- *wrk* for evaluating performance
- load test the real system